Authors: C. Amsüss, Ed.   Z. Shelby   M. Koster   C. Bormann
                                    Edge Impulse   *PassiveLogic*   *Universität Bremen TZI*

P. van der Stok
vanderstok consultancy

RFC 9176
Constrained RESTful Environments (CoRE) Resource Directory

Abstract

In many Internet of Things (IoT) applications, direct discovery of resources is not practical due to sleeping nodes or networks where multicast traffic is inefficient. These problems can be solved by employing an entity called a Resource Directory (RD), which contains information about resources held on other servers, allowing lookups to be performed for those resources. The input to an RD is composed of links, and the output is composed of links constructed from the information stored in the RD. This document specifies the web interfaces that an RD supports for web servers to discover the RD and to register, maintain, look up, and remove information on resources. Furthermore, new target attributes useful in conjunction with an RD are defined.

Status of This Memo

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in Section 2 of RFC 7841.

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at https://www.rfc-editor.org/info/rfc9176.

Copyright Notice

Table of Contents

1. Introduction

In the work on Constrained RESTful Environments (CoRE), a Representational State Transfer (REST) architecture suitable for constrained nodes (e.g., with limited RAM and ROM [RFC7228]) and networks (e.g., IPv6 over Low-Power Wireless Personal Area Network (6LoWPAN) [RFC4944]) has been established and is used in Internet of Things (IoT) or machine-to-machine (M2M) applications, such as smart energy and building automation.

The discovery of resources offered by a constrained server is very important in machine-to-machine applications where there are no humans in the loop and static interfaces result in fragility. The discovery of resources provided by an HTTP web server is typically called web linking [RFC8288]. The use of web linking for the description and discovery of resources hosted by constrained web servers is specified by the CoRE Link Format [RFC6690]. However, [RFC6690] only describes how to discover resources from the web server that hosts them by querying `/.well-known/core`. In many constrained scenarios, direct discovery of resources is not practical due to sleeping nodes or networks where multicast traffic is inefficient. These problems

can be solved by employing an entity called a Resource Directory (RD), which contains information about resources held on other servers, allowing lookups to be performed for those resources.

This document specifies the web interfaces that an RD supports for web servers to discover the RD and to register, maintain, look up, and remove information on resources. Furthermore, new target attributes useful in conjunction with an RD are defined. Although the examples in this document show the use of these interfaces with the Constrained Application Protocol (CoAP) [RFC7252], they can be applied in an equivalent manner to HTTP [RFC7230].

2. Terminology

The key words "**MUST**", "**MUST NOT**", "**REQUIRED**", "**SHALL**", "**SHALL NOT**", "**SHOULD**", "**SHOULD NOT**", "**RECOMMENDED**", "**NOT RECOMMENDED**", "**MAY**", and "**OPTIONAL**" in this document are to be interpreted as described in BCP 14 [RFC2119] [RFC8174] when, and only when, they appear in all capitals, as shown here.

The term "byte" is used in its now customary sense as a synonym for "octet".

This specification requires readers to be familiar with all the terms and concepts that are discussed in [RFC3986], [RFC8288], and [RFC6690]. Readers should also be familiar with the terms and concepts discussed in [RFC7252]. To describe the REST interfaces defined in this specification, the URI Template format is used [RFC6570].

This specification makes use of the following additional terminology:

Resolve Against
> The expression "a URI reference is *resolved against* a base URI" is used to describe the process of [RFC3986], Section 5.2. Noteworthy corner cases include the following: (1) if the URI reference is a (full) URI, resolving against any base URI gives the original full URI and (2) resolving an empty URI reference gives the base URI without any fragment identifier.

Resource Directory (RD)
> A web entity that stores information about web resources and implements the REST interfaces defined in this specification for discovery, for the creation, maintenance, and removal of registrations, and for lookup of the registered resources.

Sector
> In the context of an RD, a sector is a logical grouping of endpoints.
>
> The abbreviation "d=" is used for the sector in query parameters for compatibility with deployed implementations.

Endpoint (EP)
  Endpoint (EP) is a term used to describe a web server or client in [RFC7252]. In the context of this specification, an endpoint is used to describe a web server that registers resources to the RD. An endpoint is identified by its endpoint name, which is included during registration, and has a unique name within the associated sector of the registration.

Registration Base URI
  The base URI of a registration is a URI that typically gives scheme and authority information about an endpoint. The registration base URI is provided at registration time and is used by the RD to resolve relative references of the registration into URIs.

Target
  The target of a link is the destination address (URI) of the link. It is sometimes identified with "href=" or displayed as `<target>`. Relative targets need resolving with respect to the base URI (Section 5.2 of [RFC3986]).

  This use of the term "target" is consistent with the use in [RFC8288].

Context
  The context of a link is the source address (URI) of the link and describes which resource is linked to the target. A link's context is made explicit in serialized links as the "anchor=" attribute.

  This use of the term "context" is consistent with the use in [RFC8288].

Directory Resource
  A directory resource is a resource in the RD containing registration resources.

Registration Resource
  A registration resource is a resource in the RD that contains information about an endpoint and its links.

Commissioning Tool (CT)
  A Commissioning Tool (CT) is a device that assists during installation events by assigning values to parameters, naming endpoints and groups, or adapting the installation to the needs of the applications.

Registrant-EP
  A registrant-EP is the endpoint that is registered into the RD. The registrant-EP can register itself, or a CT registers the registrant-EP.

Resource Directory Address Option (RDAO)
  A Resource Directory Address Option (RDAO) is a new IPv6 Neighbor Discovery option defined for announcing an RD's address.

3. Architecture and Use Cases

3.1. Principles

The RD is primarily a tool to make discovery operations more efficient than querying `/.well-known/core` on all connected devices or across boundaries that would limit those operations.

It provides information about resources hosted by other devices that could otherwise only be obtained by directly querying the `/.well-known/core` resource on these other devices, either by a unicast request or a multicast request.

Information **SHOULD** only be stored in the RD if it can be obtained by querying the described device's `/.well-known/core` resource directly.

Data in the RD can only be provided by the device that hosts the data or a dedicated Commissioning Tool (CT). These CTs act on behalf of endpoints too constrained, or generally unable, to present that information themselves. No other client can modify data in the RD. Changes to the information in the RD do not propagate automatically back to the web servers from where the information originated.

3.2. Architecture

The RD architecture is illustrated in Figure 1. An RD is used as a repository of registrations describing resources hosted on other web servers, also called endpoints (EPs). An endpoint is a web server associated with a scheme, IP address, and port. A physical node may host one or more endpoints. The RD implements a set of REST interfaces for endpoints to register and maintain RD registrations and for endpoints to look up resources from the RD. An RD can be logically segmented by the use of sectors.

A mechanism to discover an RD using CoRE Link Format [RFC6690] is defined.

Registrations in the RD are soft state and need to be periodically refreshed.

An endpoint uses specific interfaces to register, update, and remove a registration. It is also possible for an RD to fetch web links from endpoints and add their contents to its registrations.

At the first registration of an endpoint, a "registration resource" is created, the location of which is returned to the registering endpoint. The registering endpoint uses this registration resource to manage the contents of registrations.

A lookup interface for discovering any of the web links stored in the RD is provided using the CoRE Link Format.

```
             Registration        Lookup
               Interface        Interface
+----+              |                |
| EP |----          |                |
+----+     ----     |                |
              --|-     +------+       |
+----+          | ----|      |       |       +--------+
| EP | ---------|-----|  RD  |----|-----| Client |
+----+          | ----|      |       |       +--------+
              --|-     +------+       |
+----+     ----     |                |
| CT |----          |                |
+----+
```
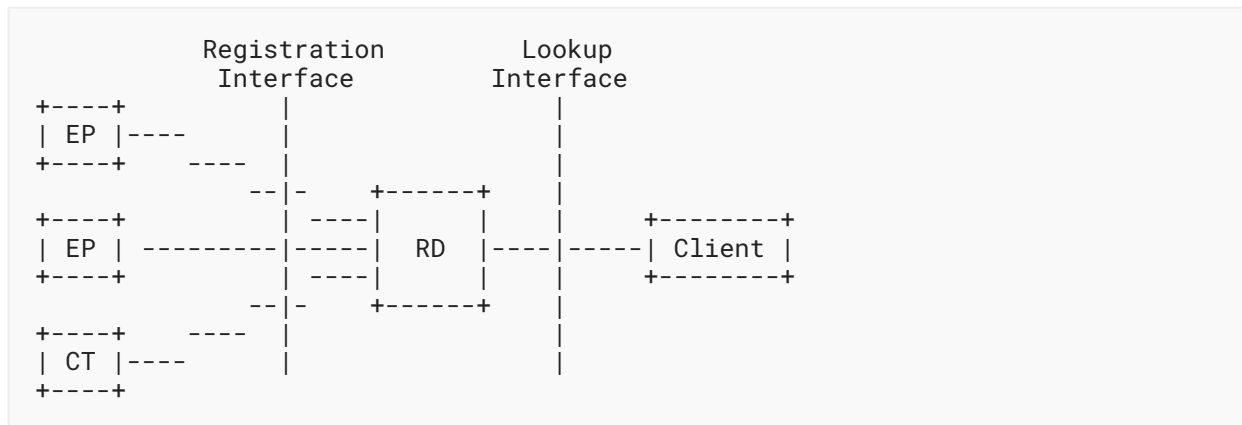
Figure 1: The RD Architecture

A registrant-EP **MAY** keep concurrent registrations to more than one RD at the same time if explicitly configured to do so, but that is not expected to be supported by typical EP implementations. Any such registrations are independent of each other. The usual expectation when multiple discovery mechanisms or addresses are configured is that they constitute a fall-back path for a single registration.

3.3. RD Content Model

The Entity-Relationship (ER) models shown in Figures 2 and 3 model the contents of `/.well-known/core` and the RD respectively, with entity-relationship diagrams [ER]. Entities (rectangles) are used for concepts that exist independently. Attributes (ovals) are used for concepts that exist only in connection with a related entity. Relations (diamonds) give a semantic meaning to the relation between entities. Numbers specify the cardinality of the relations.

Some of the attribute values are URIs. Those values are always full URIs and never relative references in the information model. However, they can be expressed as relative references in serializations, and they often are.

These models provide an abstract view of the information expressed in link-format documents and an RD. They cover the concepts but not necessarily all details of an RD's operation; they are meant to give an overview and not be a template for implementations.

```
             +----------------------+
             |   /.well-known/core  |
             +----------------------+
                        |
                        | 1
                  ////////\\\\\\\
                  <     contains   >
                  \\\\\\\\////////
                        |
                        | 0+
             +-------------------+
             |       link        |
             +-------------------+
                        |
                        |  1   oooooooo
                 +-----o target o
                 |        oooooooo
 ooooooooooooo   0+      |
 o     target  o--------+
 o  attribute o         | 0+   oooooo
  oooooooooooo           +-----o rel  o
                         |        oooooo
                         |
                         | 1    ooooooooo
                 +-----o context o
                          ooooooooo
```
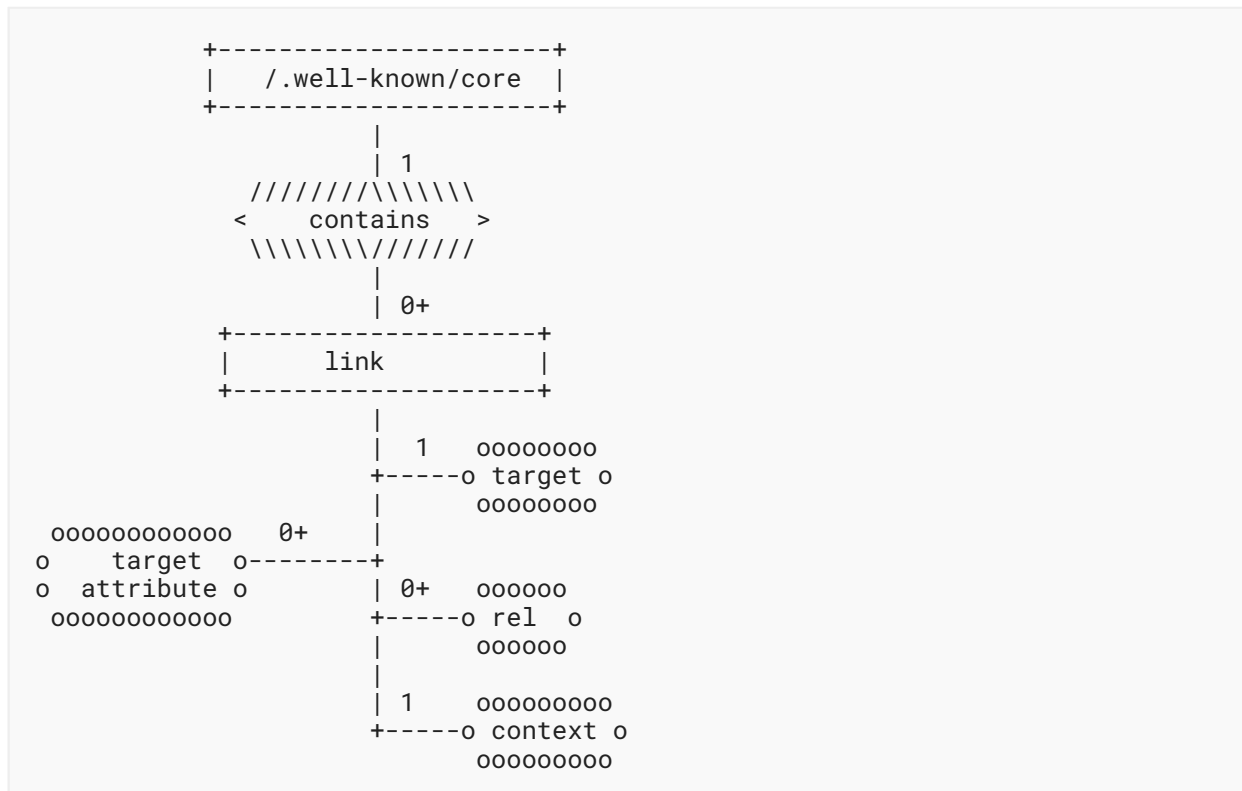
Figure 2: ER Model of the Content of `/.well-known/core`

[Figure 2](#) models the contents of `/.well-known/core`, which contains a set of links belonging to the hosting web server.

The web server is free to choose links it deems appropriate to be exposed in its `/.well-known/core`. Typically, the links describe resources that are served by the host, but the set can also contain links to resources on other servers (see examples in [Section 5](#) of [[RFC6690](#)]). The set does not necessarily contain links to all resources served by the host.

A link has the following attributes (see [Section 5](#) of [[RFC8288](#)]):

- Zero or more link relations: They describe relations between the link context and the link target.

  In link-format serialization, they are expressed as space-separated values in the "rel" attribute and default to "hosts".

- A link context URI: It defines the source of the relation, e.g., *who* "hosts" something.

  In link-format serialization, it is expressed in the "anchor" attribute and defaults to the Origin of the target (practically, the target with its path and later components removed).

- A link target URI: It defines the destination of the relation (e.g., *what* is hosted) and is the topic of all target attributes.

In link-format serialization, it is expressed between angular brackets and sometimes called
the "href".

- Other target attributes (e.g., resource type (rt), interface (if), or content format (ct)): These
provide additional information about the target URI.

```
                        +--------------+
                        +      RD      +
                        +--------------+
                               | 1
                               |
                               |
                               |
                               |
                          //////\\\\
                          < contains >
                          \\\\//////
                               |
                            0+ |
  ooooooo      1 +--------------+
 o  base o-------|  registration |
  ooooooo        +--------------+
                       |          | 1
                       |        +--------------+
    oooooooo  1 |               |
    o  href  o----+         //////\\\\
     oooooooo     |          < contains >
                  |          \\\\//////
    oooooooo  1 |               |
    o    ep   o----+              | 0+
     oooooooo     |      +------------------+
                  |      |       link       |
    oooooooo 0-1 |      +------------------+
    o    d    o----+              |
     oooooooo     |               | 1   oooooooo
                  |               +-----o target o
                  |               |     oooooooo
    oooooooo  1 |                 |
    o    lt   o----+  ooooooooooo   0+ |
     oooooooo     |  o  target   o-----+
                  |  o attribute o     | 0+   oooooo
    ooooooooooo 0+ |   ooooooooooo     +-----o rel  o
    o  endpoint o----+                 |      oooooo
    o attribute o                      |
     ooooooooooo                       | 1   oooooooo
                                       +----o context o
                                             oooooooo
```
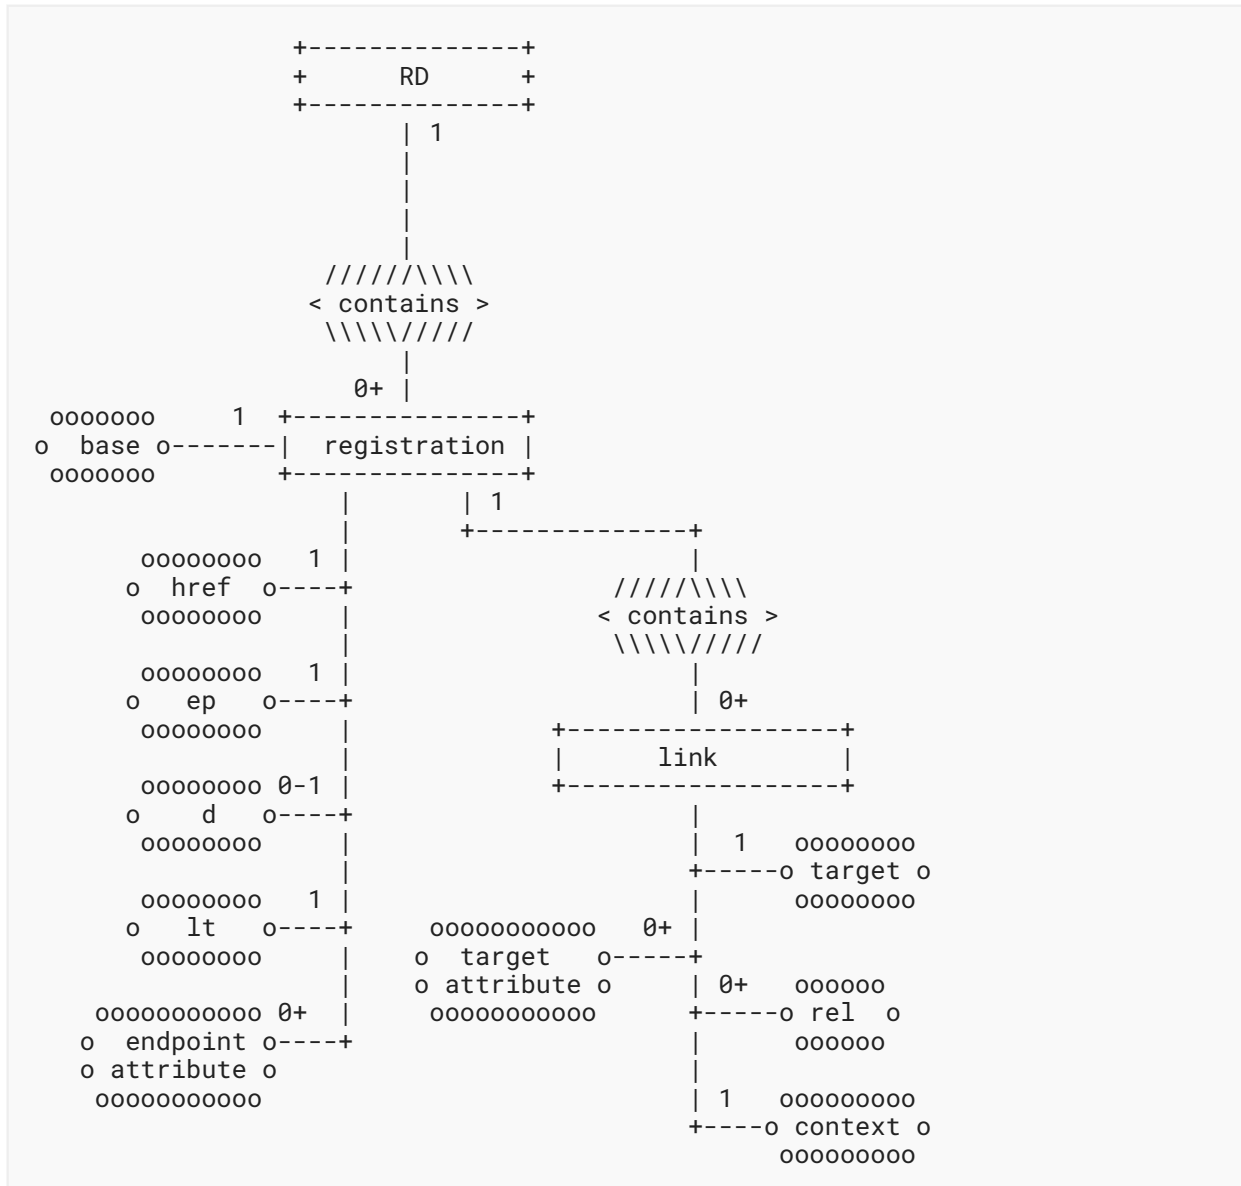
Figure 3: ER Model of the Content of the RD

Figure 3 models the contents of the RD, which contains, in addition to `/.well-known/core`, 0 to n
registrations of endpoints.

A registration is associated with one endpoint. A registration defines a set of links, as defined for `/.well-known/core`. A registration has six types of attributes:

- an endpoint name ("ep", a Unicode string) unique within a sector
- a registration base URI ("base", a URI typically describing the scheme://authority part)
- a lifetime ("lt")
- a registration resource location inside the RD ("href")
- optionally, a sector ("d", a Unicode string)
- optional additional endpoint attributes (from Section 9.3)

The cardinality of "base" is currently 1; future documents are invited to extend the RD specification to support multiple values (e.g., [COAP-PROT-NEG]). Its value is used as a base URI when resolving URIs in the links contained in the endpoint.

Links are modeled as they are in Figure 2.

3.4. Link-Local Addresses and Zone Identifiers

Registration base URIs can contain link-local IP addresses. To be usable across hosts, those cannot be serialized to contain zone identifiers (see [RFC6874], Section 1).

Link-local addresses can only be used on a single link (therefore, RD servers cannot announce them when queried on a different link), and lookup clients using them need to keep track of which interface they got them from.

Therefore, it is advisable in many scenarios to use addresses with larger scopes, if available.

3.5. Use Case: Cellular M2M

Over the last few years, mobile operators around the world have focused on development of M2M solutions in order to expand the business to the new type of users: machines. The machines are connected directly to a mobile network using an appropriate embedded wireless interface (GSM/General Packet Radio Service (GPRS), Wideband Code Division Multiple Access (W-CDMA), LTE, etc.) or via a gateway providing short- and wide-range wireless interfaces. The ambition in such systems is to build them from reusable components. These speed up development and deployment and enable shared use of machines across different applications. One crucial component of such systems is the discovery of resources (and thus the endpoints they are hosted on) capable of providing required information at a given time or acting on instructions from the end users.

Imagine a scenario where endpoints installed on vehicles enable tracking of the position of these vehicles for fleet management purposes and allow monitoring of environment parameters. During the boot-up process, endpoints register with an RD, which is hosted by the mobile operator or somewhere in the cloud. Periodically, these endpoints update their registration and may modify resources they offer.

When endpoints are not always connected, for example, because they enter a sleep mode, a remote server is usually used to provide proxy access to the endpoints. Mobile apps or web applications for environment monitoring contact the RD, look up the endpoints capable of providing information about the environment using an appropriate set of link parameters, obtain information on how to contact them (URLs of the proxy server), and then initiate interaction to obtain information that is finally processed, displayed on the screen, and usually stored in a database. Similarly, fleet management systems provide the appropriate link parameters to the RD to look up for EPs deployed on the vehicles the application is responsible for.

3.6. Use Case: Home and Building Automation

Home and commercial building automation systems can benefit from the use of IoT web services. The discovery requirements of these applications are demanding. Home automation usually relies on run-time discovery to commission the system, whereas, in building automation, a combination of professional commissioning and run-time discovery is used. Both home and building automation involve peer-to-peer interactions between endpoints and involve battery-powered sleeping devices. Both can use the common RD infrastructure to establish device interactions efficiently but can pick security policies suitable for their needs.

Two phases can be discerned for a network servicing the system: (1) installation and (2) operation. During the operational phase, the network is connected to the Internet with a border router (e.g., a 6LoWPAN Border Router (6LBR) [RFC6775]), and the nodes connected to the network can use the Internet services that are provided by the IP or network administrator. During the installation phase, the network is completely stand-alone, no border router is connected, and the network only supports the IP communication between the connected nodes. The installation phase is usually followed by the operational phase. As an RD's operations work without hard dependencies on names or addresses, it can be used for discovery across both phases.

3.7. Use Case: Link Catalogues

Resources may be shared through data brokers that have no knowledge beforehand of who is going to consume the data. An RD can be used to hold links about resources and services hosted anywhere to make them discoverable by a general class of applications.

For example, environmental and weather sensors that generate data for public consumption may provide data to an intermediary server or broker. Sensor data are published to the intermediary upon changes or at regular intervals. Descriptions of the sensors that resolve to links to sensor data may be published to an RD. Applications wishing to consume the data can use RD lookup to discover and resolve links to the desired resources and endpoints. The RD service need not be coupled with the data intermediary service. Mapping of RDs to data intermediaries may be many-to-many.

Metadata in web link formats, such as the one defined in [RFC6690], which may be internally stored as triples or relation/attribute pairs providing metadata about resource links, need to be supported by RDs. External catalogues that are represented in other formats may be converted to

common web linking formats for storage and access by RDs. Since it is common practice for these to be encoded in URNs [RFC8141], simple and lossless structural transforms should generally be sufficient to store external metadata in RDs.

The additional features of an RD allow sectors to be defined to enable access to a particular set of resources from particular applications. This provides isolation and protection of sensitive data when needed. Application groups with multicast addresses may be defined to support efficient data transport.

4. RD Discovery and Other Interface-Independent Components

This and the following sections define the required set of REST interfaces between an RD, endpoints, and lookup clients. Although the examples throughout these sections assume the use of CoAP [RFC7252], these REST interfaces can also be realized using HTTP [RFC7230]. The multicast discovery and simple registration operations are exceptions to that, as they rely on mechanisms unavailable in HTTP. In all definitions in these sections, both CoAP response codes (with dot notation) and HTTP response codes (without dot notation) are shown. An RD implementing this specification **MUST** support the discovery, registration, update, lookup, and removal interfaces.

All operations on the contents of the RD **MUST** be atomic and idempotent.

For several operations, interface templates are given in list form; those describe the operation participants, request codes, URIs, content formats, and outcomes. Sections of those templates contain normative content about Interaction, Method, URI Template, and URI Template Variables, as well as the details of the Success condition. The additional sections for options (such as Content-Format) and for Failure codes give typical cases that an implementation of the RD should deal with. Those serve to illustrate the typical responses to readers who are not yet familiar with all the details of CoAP-based interfaces; they do not limit how a server may respond under atypical circumstances.

REST clients (registrant-EPs and CTs during registration and maintenance, lookup clients, and RD servers during simple registrations) must be prepared to receive any unsuccessful code and act upon it according to its definition, options, and/or payload to the best of their capabilities, falling back to failing the operation if recovery is not possible. In particular, they **SHOULD** retry the request upon 5.03 (Service Unavailable; 503 in HTTP) according to the Max-Age (Retry-After in HTTP) option and **SHOULD** fall back to link format when receiving 4.15 (Unsupported Content-Format; 415 in HTTP).

An RD **MAY** make the information submitted to it available to further directories (subject to security policies on link confidentiality) if it can ensure that a loop does not form. The protocol used between directories to ensure loop-free operation is outside the scope of this document.

4.1. Finding a Resource Directory

A (re)starting device may want to find one or more RDs before it can discover their URIs. Dependent on the operational conditions, one or more of the techniques below apply.

The device may be preconfigured to exercise specific mechanisms for finding the RD:

1. It may be configured with a specific IP address for the RD. That IP address may also be an anycast address, allowing the network to forward RD requests to an RD that is topologically close; each target network environment in which some of these preconfigured nodes are to be brought up is then configured with a route for this anycast address that leads to an appropriate RD. (Instead of using an anycast address, a multicast address can also be preconfigured. The RD servers then need to configure one of their interfaces with this multicast address.)

2. It may be configured with a DNS name for the RD and use DNS to return the IP address of the RD; it can find a DNS server to perform the lookup using the usual mechanisms for finding DNS servers.

3. It may be configured to use a service discovery mechanism, such as DNS-based Service Discovery (DNS-SD), as outlined in Section 4.1.2.

For cases where the device is not specifically configured with a way to find an RD, the network may want to provide a suitable default.

1. The IPv6 Neighbor Discovery option RDAO (Section 4.1.1) can do that.

2. When DHCP is in use, this could be provided via a DHCP option (no such option is defined at the time of writing).

Finally, if neither the device nor the network offers any specific configuration, the device may want to employ heuristics to find a suitable RD.

The present specification does not fully define these heuristics but suggests a number of candidates:

1. In a 6LoWPAN, just assume the 6LBR can act as an RD (using the Authoritative Border Router Option (ABRO) to find that [RFC6775]). Confirmation can be obtained by sending a unicast GET to `coap://[6LBR]/.well-known/core?rt=core.rd*`.

2. In a network that supports multicast well, discover the RD using a multicast query for `/.well-known/core`, as specified in CoRE Link Format [RFC6690], and send a Multicast GET to `coap://[ff0x::fe]/.well-known/core?rt=core.rd*`. RDs within the multicast scope will answer the query.

When answering a multicast request directed at a link-local group, the RD may want to respond from a routable address; this makes it easier for registrants to use one of their own routable addresses for registration. When source addresses are selected using the mechanism described in [RFC6724], this can be achieved by applying the changes of its Section 10.4, picking public addresses in Rule 7 of its Section 5, and superseding Rule 8 with preferring the source address's precedence.

As some of the RD addresses obtained by the methods listed here are just (more or less educated) guesses, endpoints **MUST** make use of any error messages to very strictly rate-limit requests to candidate IP addresses that don't work out. For example, an ICMP Destination Unreachable

message (and, in particular, the port unreachable code for this message) may indicate the lack of a CoAP server on the candidate host, or a CoAP error response code, such as 4.05 (Method Not Allowed), may indicate unwillingness of a CoAP server to act as a directory server.

The following RD discovery mechanisms are recommended:

- In managed networks with border routers that need stand-alone operation, the RDAO is recommended (e.g., the operational phase described in Section 3.6).
- In managed networks without border routers (no Internet services available), the use of a preconfigured anycast address is recommended (e.g., the installation phase described in Section 3.6).
- In networks managed using DNS-SD, the use of DNS-SD for discovery, as described in Section 4.1.2, is recommended.

The use of multicast discovery in mesh networks is **NOT RECOMMENDED**.

4.1.1. Resource Directory Address Option (RDAO)

The Resource Directory Address Option (RDAO) carries information about the address of the RD in RAs (Router Advertisements) of IPv6 Neighbor Discovery (ND), similar to how Recursive DNS Server (RDNSS) options [RFC8106] are sent. This information is needed when endpoints cannot discover the RD with a link-local or realm-local scope multicast address, for instance, because the endpoint and the RD are separated by a 6LBR. In many circumstances, the availability of DHCP cannot be guaranteed during commissioning of the network either. The presence and the use of the RD is essential during commissioning.

It is possible to send multiple RDAOs in one message, indicating as many RD addresses.
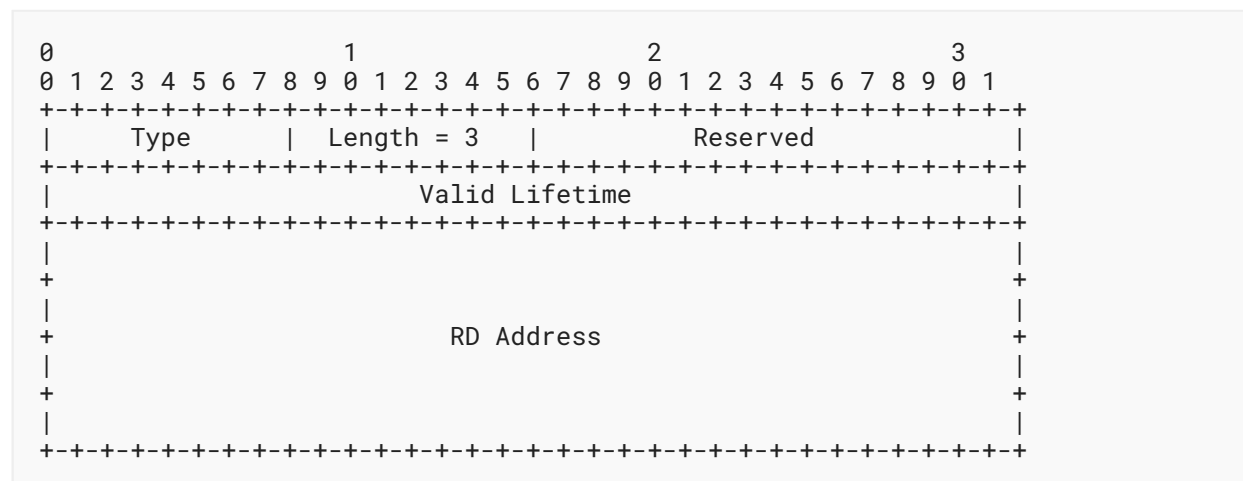
The RDAO format is:

```
 0                   1                   2                   3
 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1 2 3 4 5 6 7 8 9 0 1
+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+
|     Type      |  Length = 3   |            Reserved           |
+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+
|                         Valid Lifetime                        |
+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+
|                                                               |
+                                                               +
|                                                               |
+                         RD Address                            +
|                                                               |
+                                                               +
|                                                               |
+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+-+
```

Figure 4: Resource Directory Address Option

Fields:

Type: 41

Length: 8-bit unsigned integer. The length of the option in units of 8 bytes. Always 3.

Reserved: This field is unused. It **MUST** be initialized to zero by the sender and **MUST** be ignored by the receiver.

Valid Lifetime: 32-bit unsigned integer. The length of time in seconds (relative to the time the packet is received) that this RD address is valid. A value of all zero bits (0x0) indicates that this RD address is not valid anymore.

RD Address: IPv6 address of the RD.

4.1.2. Using DNS-SD to Discover a Resource Directory

An RD can advertise its presence in DNS-SD [RFC6763] using the service names defined in this document: `_core-rd._udp` (for CoAP), `_core-rd-dtls._udp` (for CoAP over DTLS), `_core-rd._tcp` (for CoAP over TCP), or `_core-rd-tls._tcp` (for CoAP over TLS). (For the WebSocket transports of CoAP, no service is defined, as DNS-SD is typically unavailable in environments where CoAP over WebSockets is used.)

The selection of the service indicates the protocol used, and the SRV record points the client to a host name and port to use as a starting point for the "URI discovery" steps of Section 4.3.

This section is a simplified, concrete application of the more generic mechanism specified in [CORE-RD-DNS-SD].

4.2. Payload Content Formats

RDs implementing this specification **MUST** support the `application/link-format` content format (ct=40).

RDs implementing this specification **MAY** support additional content formats.

Any additional content format supported by an RD implementing this specification **SHOULD** be able to express all the information expressible in link format. It **MAY** be able to express information that is inexpressible in link format, but those expressions **SHOULD** be avoided where possible.

4.3. URI Discovery

Before an endpoint can make use of an RD, it must first know the RD's address and port and the URI path information for its REST APIs. This section defines discovery of the RD and its URIs using the well-known interface of the CoRE Link Format [RFC6690] after having discovered a host, as described in Section 4.1.

Discovery of the RD registration URI is performed by sending either a multicast or unicast GET request to `/.well-known/core` and including a resource type (rt) parameter [RFC6690] with the value "core.rd" in the query string. Likewise, a resource type parameter value of "core.rd-

lookup*" is used to discover the URIs for RD lookup operations, and "core.rd*" is used to discover all URIs for RD operations. Upon success, the response will contain a payload with a link format entry for each RD function discovered, indicating the URI of the RD function returned and the corresponding resource type. When performing multicast discovery, the multicast IP address used will depend on the scope required and the multicast capabilities of the network (see Section 9.5).

An RD **MAY** provide hints about the content formats it supports in the links it exposes or registers, using the "ct" target attribute, as shown in the example below. Clients **MAY** use these hints to select alternate content formats for interaction with the RD.

HTTP does not support multicast, and, consequently, only unicast discovery can be supported using the HTTP `/.well-known/core` resource.

RDs implementing this specification **MUST** support query filtering for the rt parameter, as defined in [RFC6690].

While the link targets in this discovery step are often expressed in path-absolute form, this is not a requirement. Clients of the RD **SHOULD** therefore accept URIs of all schemes they support, both as URIs and relative references, and not limit the set of discovered URIs to those hosted at the address used for URI discovery.

With security policies where the client requires the RD to be authorized to act as an RD, that authorization may be limited to resources on which the authorized RD advertises the adequate resource types. Clients that have obtained links they cannot rely on yet can repeat the "URI discovery" step at the `/.well-known/core` resource of the indicated host to obtain the resource type information from an authorized source.

The URI discovery operation can yield multiple URIs of a given resource type. The client of the RD can try out any of the discovered addresses.

The discovery request interface is specified as follows (this is exactly the well-known interface of [RFC6690], Section 4, with the additional requirement that the server **MUST** support query filtering):

Interaction: EP, CT, or Client -> RD

Method: GET

URI Template: `/.well-known/core{?rt}`

URI Template Variables:

   rt :=   Resource Type. **SHOULD** contain one of the values "core.rd", "core.rd-lookup*", "core.rd-lookup-res", "core.rd-lookup-ep", or "core.rd*"

Accept: absent, `application/link-format`, or any other media type representing web links

The following response is expected on this interface:

Success: 2.05 (Content) or 200 (OK) with an `application/link-format` or other web link
payload containing one or more matching entries for the RD resource.

The following example shows an endpoint discovering an RD using this interface, thus learning
that the directory resource location in this example is /rd and that the content format delivered
by the server hosting the resource is `application/link-format` (ct=40). Note that it is up to the
RD to choose its RD locations.

```
Req: GET coap://[ff02::fe]/.well-known/core?rt=core.rd*

Res: 2.05 Content
Payload:
</rd>;rt=core.rd;ct=40,
</rd-lookup/ep>;rt=core.rd-lookup-ep;ct=40,
</rd-lookup/res>;rt=core.rd-lookup-res;ct=40
```

Figure 5: Example Discovery Exchange

The following example shows the way of indicating that a client may request alternate content
formats. The Content-Format code attribute "ct" **MAY** include a space-separated sequence of
Content-Format codes, as specified in Section 7.2.1 of [RFC7252], indicating that multiple content
formats are available. The example below shows the required Content-Format 40 (`application/
link-format`) indicated, as well as Concise Binary Object Representation (CBOR) and JSON
representations in the style of [CORE-LINKS-JSON] (for which the experimental values 65060 and
65050 are used in this example). The RD resource locations /rd and /rd-lookup are example
values. The server in this example also indicates that it is capable of providing observation on
resource lookups.

```
Req: GET coap://[ff02::fe]/.well-known/core?rt=core.rd*

Res: 2.05 Content
Payload:
</rd>;rt=core.rd;ct=40,
</rd-lookup/res>;rt=core.rd-lookup-res;ct="40 65060 65050";obs,
</rd-lookup/ep>;rt=core.rd-lookup-ep;ct="40 65060 65050"
```

Figure 6: Example Discovery Exchange Indicating Additional Content-Formats

For maintenance, management, and debugging, it can be useful to identify the components that
constitute the RD server. The identification can be used to find client-server incompatibilities,
supported features, required updates, and other aspects. The well-known interface described in
Section 4 of [RFC6690] can be used to find such data.

It would typically be stored in an implementation information link (as described in [T2TRG-REL-
IMPL]).

```
Req: GET /.well-known/core?rel=impl-info

Res: 2.05 Content
Payload:
<http://software.example.com/shiny-resource-directory/1.0beta1>;
    rel=impl-info
```

Figure 7: Example Exchange of Obtaining Implementation Information Using the Relation Type
Currently Proposed in [T2TRG-REL-IMPL]

Note that, depending on the particular server's architecture, such a link could be anchored at the
RD server's root (as in this example) or at individual RD components. The latter is to be expected
when different applications are run on the same server.

5. Registration

After discovering the location of an RD, a registrant-EP or CT **MAY** register the resources of the
registrant-EP using the registration interface. This interface accepts a POST from an endpoint
containing the list of resources to be added to the directory as the message payload in the CoRE
Link Format [RFC6690] or other representations of web links, along with query parameters
indicating the name of the endpoint and, optionally, the sector, lifetime, and base URI of the
registration. It is expected that other specifications will define further parameters (see Section
9.3). The RD then creates a new registration resource in the RD and returns its location. The
receiving endpoint **MUST** use that location when refreshing registrations using this interface.
Registration resources in the RD are kept active for the period indicated by the lifetime
parameter. The creating endpoint is responsible for refreshing the registration resource within
this period, using either the registration or update interface. The registration interface **MUST** be
implemented to be idempotent, so that registering twice with the same endpoint parameters ep
and d (sector) does not create multiple registration resources.

The following rules apply for a registration request targeting a given (ep, d) value pair:

- When the (ep, d) value pair of the registration request is different from any existing
  registration, a new registration is generated.
- When the (ep, d) value pair of the registration request is equal to an existing registration, the
  content and parameters of the existing registration are replaced with the content of the
  registration request. As with changes to registration resources, security policies (Section 7)
  usually require such requests to come from the same device.

The posted link-format document can (and typically does) contain relative references both in its
link targets and in its anchors; it can also contain empty anchors. The RD server needs to resolve
these references in order to faithfully represent them in lookups. They are resolved against the
base URI of the registration, which is provided either explicitly in the `base` parameter or
constructed implicitly from the requester's URI, as constructed from its network address and
scheme.

For media types to which Appendix C applies (i.e., documents in `application/link-format`), request bodies **MUST** be expressed in Limited Link Format.

The registration request interface is specified as follows:

Interaction: EP or CT -> RD

Method: POST

URI Template: {+rd}{?ep,d,lt,base,extra-attrs*}

URI Template Variables:

rd :=   RD registration URI (mandatory). This is the location of the RD, as obtained from discovery.

ep :=   Endpoint name (mostly mandatory). The endpoint name is an identifier that **MUST** be unique within a sector.

As the endpoint name is a Unicode string, it is encoded in UTF-8 (and possibly percent encoded) during variable expansion (see [RFC6570], Section 3.2.1). The endpoint name **MUST NOT** contain any character in the inclusive ranges 0-31 or 127-159.

The maximum length of this parameter is 63 bytes encoded in UTF-8.

If the RD is configured to recognize the endpoint that is to be authorized to use exactly one endpoint name, the RD assigns that name. In that case, giving the endpoint name becomes optional for the client; if the client gives any other endpoint name, it is not authorized to perform the registration.

d :=   Sector (optional). This is the sector to which this endpoint belongs. When this parameter is not present, the RD **MAY** associate the endpoint with a configured default sector (possibly based on the endpoint's authorization) or leave it empty.

The sector is encoded like the ep parameter and is limited to 63 bytes encoded in UTF-8 as well.

lt :=   Lifetime (optional). This is the lifetime of the registration in seconds, with a range of 1-4294967295. If no lifetime is included in the initial registration, a default value of 90000 (25 hours) **SHOULD** be assumed.

base :=   Base URI (optional). This parameter sets the base URI of the registration, under which the relative links in the payload are to be interpreted. The specified URI typically does not have a path component of its own and **MUST** be suitable as a base URI to resolve any relative references given in the registration. The parameter is therefore usually of the shape "scheme://authority" for HTTP and CoAP URIs. The URI **SHOULD NOT** have a query or fragment component, as any non-empty relative part in a reference would remove those parts from the resulting URI.

In the absence of this parameter, the scheme of the protocol, the source address, and the source port of the registration request are assumed. The base URI is consecutively constructed by concatenating the used protocol's scheme with the characters "://", the requester's source address as an address literal, and ":" followed by its port (if it was not the protocol's default one). This is analogous to the process described in [RFC7252], Section 6.5.

This parameter is mandatory when the directory is filled by a third party, such as a commissioning tool.

If the registrant-EP uses an ephemeral port to register with, it **MUST** include the base parameter in the registration to provide a valid network path.

A registrant that cannot be reached by potential lookup clients at the address it registers from (e.g., because it is behind some form of Network Address Translation (NAT)) **MUST** provide a reachable base address with its registration.

If the base URI contains a link-local IP literal, it **MUST NOT** contain a Zone Identifier and **MUST** be local to the link on which the registration request is received.

Endpoints that register with a base that contains a path component cannot efficiently express their registrations in Limited Link Format (Appendix C). Those applications should use different representations of links to which Appendix C is not applicable (e.g., [CORE-CORAL]).

extra-attrs :=   Additional registration attributes (optional). The endpoint can pass any parameter registered in Section 9.3 to the directory. If the RD is aware of the parameter's specified semantics, it processes the parameter accordingly. Otherwise, it **MUST** store the unknown key and its value(s) as an endpoint attribute for further lookup.

Content-Format: `application/link-format` or any other indicated media type representing web links

The following response is expected on this interface:

Success: 2.01 (Created) or 201 (Created). The Location-Path option or Location header field **MUST** be included in the response. This location **MUST** be a stable identifier generated by the RD, as it is used for all subsequent operations on this registration resource. The registration resource location thus returned is for the purpose of updating the lifetime of the registration and for maintaining the content of the registered links, including updating and deleting links.

A registration with an already-registered ep and d value pair responds with the same success code and location as the original registration; the set of links registered with the endpoint is replaced with the links from the payload.

The location **MUST NOT** have a query or fragment component, as that could conflict with query parameters during the registration update operation. Therefore, the Location-Query option **MUST NOT** be present in a successful response.

If the registration fails, including request timeouts, or if delays from Service Unavailable responses with Max-Age or Retry-After accumulate to exceed the registrant's configured timeouts, it **SHOULD** pick another registration URI from the "URI discovery" step of Section 4.3, and, if there is only one or the list is exhausted, pick other choices from the "finding a resource directory" step of Section 4.1. Care has to be taken to consider the freshness of results obtained earlier, e.g., the result of a `/.well-known/core` response, the lifetime of an RDAO, and DNS responses. Any rate limits and persistent errors from the "finding a resource directory" step must be considered for the whole registration time, not only for a single operation.

The following example shows a registrant-EP with the name "node1" registering two resources to an RD using this interface. The location "/rd" is an example RD location discovered in a request similar to Figure 5.

```
Req: POST coap://rd.example.com/rd?ep=node1
Content-Format: 40
Payload:
</sensors/temp>;rt=temperature-c;if=sensor,
<http://www.example.com/sensors/temp>;
  anchor="/sensors/temp";rel=describedby

Res: 2.01 Created
Location-Path: /rd/4521
```

Figure 8: Example Registration Payload

An RD may optionally support HTTP. Here is an example of almost the same registration operation above when done using HTTP.

```
Req:
POST /rd?ep=node1&base=http://[2001:db8:1::1] HTTP/1.1
Host: rd.example.com
Content-Type: application/link-format

</sensors/temp>;rt=temperature-c;if=sensor,
<http://www.example.com/sensors/temp>;
  anchor="/sensors/temp";rel=describedby

Res:
HTTP/1.1 201 Created
Location: /rd/4521
```

Figure 9: Example Registration Payload as Expressed Using HTTP

5.1. Simple Registration

Not all endpoints hosting resources are expected to know how to upload links to an RD, as described in Section 5. Instead, simple endpoints can implement the simple registration approach described in this section. An RD implementing this specification **MUST** implement simple registration. However, there may be security reasons why this form of directory discovery would be disabled.

This approach requires that the registrant-EP makes available the hosted resources that it wants to be discovered as links on its `/.well-known/core` interface, as specified in [RFC6690]. The links in that document are subject to the same limitations as the payload of a registration (with respect to Appendix C).

- The registrant-EP finds one or more addresses of the directory server, as described in Section 4.1.
- The registrant-EP sends (and regularly refreshes with) a POST request to the `/.well-known/rd` URI of the directory server of choice. The body of the POST request is empty and triggers the resource directory server to perform GET requests (redone before lifetime expiry) at the requesting registrant-EP's `/.well-known/core` to obtain the link-format payload to register.

  The registrant-EP includes the same registration parameters in the POST request as it would with a regular registration, per Section 5. The registration base URI of the registration is taken from the registrant-EP's network address (as is default with regular registrations).

  The following is an example request from the registrant-EP to the RD (unanswered until the next step):

  ```
  Req: POST /.well-known/rd?lt=6000&ep=node1
  (No payload)
  ```

Figure 10: First-Half Example Exchange of a Simple Registration

- The RD queries the registrant-EP's discovery resource to determine the success of the operation. It **SHOULD** keep a cache of the discovery resource and not query it again as long as it is fresh.

  The following is an example request from the RD to the registrant-EP:

  ```
  Req: GET /.well-known/core
  Accept: 40

  Res: 2.05 Content
  Content-Format: 40
  Payload:
  </sen/temp>
  ```

Figure 11: Example Exchange of the RD Querying the Simple Endpoint

With this response, the RD would answer the previous step's request:

```
Res: 2.04 Changed
```

Figure 12: Second-Half Example Exchange of a Simple Registration

The sequence of fetching the registration content before sending a successful response was chosen to make responses reliable, and the point about caching was chosen to still allow very constrained registrants. Registrants **MUST** be able to serve a GET request to `/.well-known/core` after having requested registration. Constrained devices **MAY** regard the initial request as temporarily failed when they need RAM occupied by their own request to serve the RD's GET and retry later when the RD already has a cached representation of their discovery resources. Then, the RD can reply immediately, and the registrant can receive the response.

The simple registration request interface is specified as follows:

Interaction: EP -> RD

Method: POST

URI Template: `/.well-known/rd{?ep,d,lt,extra-attrs*}`

URI Template Variables are the same as for registration in Section 5. The base attribute is not accepted to keep the registration interface simple; that rules out registration over CoAP-over-TCP or HTTP that would need to specify one. For some time during this document's development, the URI Template `/.well-known/core{?ep,...}` was in use instead.

The following response is expected on this interface:

Success: 2.04 (Changed)

For the second interaction triggered by the above, the registrant-EP takes the role of server and the RD takes the role of client. (Note that this is exactly the well-known interface of [RFC6690], Section 4):

Interaction: RD -> EP

Method: GET

URI Template: `/.well-known/core`

The following response is expected on this interface:

Success: 2.05 (Content)

When the RD uses any authorization credentials to access the endpoint's discovery resource or when it is deployed in a location where third parties might reach it but not the endpoint, it **SHOULD** verify that the apparent registrant-EP intends to register with the given registration parameters before revealing the obtained discovery information to lookup clients. An easy way to do that is to verify the simple registration request's sender address using the Echo option, as described in [RFC9175], Section 2.4.

The RD **MUST** delete registrations created by simple registration after the expiration of their lifetime. Additional operations on the registration resource cannot be executed because no registration location is returned.

5.2. Third-Party Registration

For some applications, even simple registration may be too taxing for some very constrained devices, in particular, if the security requirements become too onerous.

In a controlled environment (e.g., building control), the RD can be filled by a third-party device, called a Commissioning Tool (CT). The CT can fill the RD from a database or other means. For that purpose scheme, the IP address and port of the URI of the registered device is the value of the "base" parameter of the registration described in Section 5.

It should be noted that the value of the "base" parameter applies to all the links of the registration and has consequences for the anchor value of the individual links, as exemplified in Appendix B. A potential (currently nonexistent) "base" attribute of the link is not affected by the value of "base" parameter in the registration.

5.3. Operations on the Registration Resource

This section describes how the registering endpoint can maintain the registrations that it created. The registering endpoint can be the registrant-EP or the CT. The registrations are resources of the RD.

An endpoint should not use this interface for registrations that it did not create. This is usually enforced by security policies, which, in general, require equivalent credentials for creation of and operations on a registration.

After the initial registration, the registering endpoint retains the returned location of the registration resource for further operations, including refreshing the registration in order to extend the lifetime and "keep-alive" the registration. When the lifetime of the registration has expired, the RD **SHOULD NOT** respond to discovery queries concerning this endpoint. The RD **SHOULD** continue to provide access to the registration resource after a registration timeout occurs in order to enable the registering endpoint to eventually refresh the registration. The RD **MAY** eventually remove the registration resource for the purpose of garbage collection. If the registration resource is removed, the corresponding endpoint will need to be reregistered.

The registration resource may also be used to cancel the registration using DELETE and to perform further operations beyond the scope of this specification.

Operations on the registration resource are sensitive to reordering; Section 5.3.4 describes how order is restored.

The operations on the registration resource are described below.

5.3.1. Registration Update

The update interface is used by the registering endpoint to refresh or update its registration with an RD. To use the interface, the registering endpoint sends a POST request to the registration resource returned by the initial registration operation.

An update **MAY** update registration parameters, such as lifetime, base URI, or others. Parameters that are not being changed should not be included in an update. Adding parameters that have not changed increases the size of the message but does not have any other implications. Parameters are included as query parameters in an update operation, as in Section 5.

A registration update resets the timeout of the registration to the (possibly updated) lifetime of the registration, independent of whether an `lt` parameter was given.

If the base URI of the registration is changed in an update, relative references submitted in the original registration or later updates are resolved anew against the new base.

The registration update operation only describes the use of POST with an empty payload. Future standards might describe the semantics of using content formats and payloads with the POST method to update the links of a registration (see Section 5.3.3).

The update registration request interface is specified as follows:

Interaction: EP or CT -> RD

Method: POST

URI Template: {+location}{?lt,base,extra-attrs*}

URI Template Variables:

location :=   This is the location returned by the RD as a result of a successful earlier registration.

lt :=   Lifetime (optional). This is the lifetime of the registration in seconds, with a range of 1-4294967295. If no lifetime is included, the previous last lifetime set on a previous update or the original registration (falling back to 90000) **SHOULD** be used.

base :=   Base URI (optional). This parameter updates the base URI established in the original registration to a new value and is subject to the same restrictions as in the registration.

If the parameter is set in an update, it is stored by the RD as the new base URI under which to interpret the relative links present in the payload of the original registration.

> If the parameter is not set in the request but was set before, the previous base URI value is kept unmodified.
>
> If the parameter is not set in the request and was not set before either, the source address and source port of the update request are stored as the base URI.

extra-attrs :=   Additional registration attributes (optional). As with the registration, the RD processes them if it knows their semantics. Otherwise, unknown attributes are stored as endpoint attributes, overriding any previously stored endpoint attributes of the same key.

> Note that this default behavior does not allow removing an endpoint attribute in an update. For attributes whose functionality depends on the endpoints' ability to remove them in an update, it can make sense to define a value whose presence is equivalent to the absence of a value. As an alternative, an extension can define different updating rules for their attributes. That necessitates either discovering whether the RD is aware of that extension or tolerating the default behavior.

Content-Format: none (no payload)

The following responses are expected on this interface:

Success: 2.04 (Changed) or 204 (No Content) if the update was successfully processed.

Failure: 4.04 (Not Found) or 404 (Not Found). Registration does not exist (e.g., may have been removed).

If the registration update fails in any way, including "Not Found" and request timeouts, or if the time indicated in a Service Unavailable Max-Age/Retry-After exceeds the remaining lifetime, the registering endpoint **SHOULD** attempt registration again.

The following example shows how the registering endpoint resets the timeout on its registration resource at an RD using this interface with the example location value /rd/4521:

```
Req: POST /rd/4521

Res: 2.04 Changed
```

Figure 13: Example Update of a Registration

The following example shows the registering endpoint updating its registration resource at an RD using this interface with the example location value /rd/4521. The initial registration by the registering endpoint set the following values:

- endpoint name (ep)=endpoint1
- lifetime (lt)=500
- base URI (base)=coap://local-proxy-old.example.com
- payload of Figure 8

The initial state of the RD is reflected in the following request:

```
Req: GET /rd-lookup/res?ep=endpoint1

Res: 2.05 Content
Payload:
<coap://local-proxy-old.example.com/sensors/temp>;
    rt=temperature-c;if=sensor,
<http://www.example.com/sensors/temp>;
    anchor="coap://local-proxy-old.example.com/sensors/temp";
    rel=describedby
```

Figure 14: Example Lookup Before a Change to the Base Address

The following example shows the registering endpoint changing the base URI to `coaps://new.example.com:5684`:

```
Req: POST /rd/4521?base=coaps://new.example.com

Res: 2.04 Changed
```

Figure 15: Example Registration Update that Changes the Base Address

The consecutive query returns:

```
Req: GET /rd-lookup/res?ep=endpoint1

Res: 2.05 Content
Payload:
<coaps://new.example.com/sensors/temp>;
    rt=temperature-c;if=sensor,
<http://www.example.com/sensors/temp>;
    anchor="coaps://new.example.com/sensors/temp";
    rel=describedby
```

Figure 16: Example Lookup After a Change to the Base Address

5.3.2. Registration Removal

Although RD registrations have soft state and will eventually time out after their lifetime, the registering endpoint **SHOULD** explicitly remove an entry from the RD if it knows it will no longer be available (for example, on shutdown). This is accomplished using a removal interface on the RD by performing a DELETE on the endpoint resource.

The removal request interface is specified as follows:

Interaction: EP or CT -> RD

Method: DELETE

URI Template: {+location}

URI Template Variables:

   location :=   This is the location returned by the RD as a result of a successful earlier
        registration.

The following responses are expected on this interface:

Success: 2.02 (Deleted) or 204 (No Content) upon successful deletion.

Failure: 4.04 (Not Found) or 404 (Not Found). Registration does not exist (e.g., may already have
        been removed).

The following example shows successful removal of the endpoint from the RD with example
location value /rd/4521:

```
Req: DELETE /rd/4521

Res: 2.02 Deleted
```

Figure 17: Example of a Registration Removal

5.3.3. Further Operations

Additional operations on the registration can be specified in future documents, for example:

- Send iPATCH (or PATCH) updates [RFC8132] to add, remove, or change the links of a
  registration.
- Use GET to read the currently stored set of links in a registration resource.

Those operations are out of scope of this document and will require media types suitable for
modifying sets of links.

5.3.4. Request Freshness

Some security mechanisms usable with an RD allow out-of-order request processing or do not
even mandate replay protection at all. The RD needs to ensure that operations on the registration
resource are executed in an order that does not distort the client's intentions.

This ordering of operations is expressed in terms of freshness, as defined in [RFC9175]. Requests
that alter a resource's state need to be fresh relative to the latest request that altered that state in
a conflicting way.

An RD **SHOULD** determine a request's freshness and **MUST** use the Echo option if it requires
request freshness and cannot determine it in any other way. An endpoint **MUST** support the use
of the Echo option. (One reason why an RD would not require freshness is when no relevant
registration properties are covered by its security policies.)

5.3.4.1. Efficient Use of Echo by an RD

To keep latency and traffic added by the freshness requirements to a minimum, RDs should avoid naive (sufficient but inefficient) freshness criteria.

Some simple mechanisms the RD can employ are:

- State counter. The RD can keep a monotonous counter that increments whenever a registration changes. For every registration resource, it stores the post-increment value of that resource's last change. Requests altering them need to have at least that value encoded in their Echo option and are otherwise rejected with a 4.01 (Unauthorized) and the current counter value as the Echo value. If other applications on the same server use Echo as well, that encoding may include a prefix indicating that it pertains to the RD's counter.

  The value associated with a resource needs to be kept across the removal of registrations if the same registration resource is to be reused.

  The counter can be reset (and the values of removed resources forgotten) when all previous security associations are reset.

  This is the "Persistent Counter" method of [RFC9175], Appendix A.

- Preemptive Echo values. The current state counter can be sent in an Echo option not only when requests are rejected with 4.01 (Unauthorized) but also with successful responses. Thus, clients can be provided with Echo values sufficient for their next request on a regular basis. This is also described in Section 2.3 of [RFC9175]

  While endpoints may discard received Echo values at leisure between requests, they are encouraged to retain these values for the next request to avoid additional round trips.

- If the RD can ensure that only one security association has modifying access to any registration at any given time and that security association provides order on the requests, that order is sufficient to show request freshness.

5.3.4.2. Examples of Echo Usage

Figure 18 shows the interactions of an endpoint that has forgotten the server's latest Echo value and temporarily reduces its registration lifetime:

```
Req: POST /rd/4521?lt=7200

Res: 4.01 Unauthorized
Echo: 0x0123

(EP tries again immediately.)

Req: POST /rd/4521?lt=7200
Echo: 0x0123

Res: 2.04 Changed
Echo: 0x0124

(Later, the EP regains its confidence in its long-term reachability.)

Req: POST /rd/4521?lt=90000
Echo: 0x0124

Res: 2.04 Changed
Echo: 0x0247
```

Figure 18: Example Update of a Registration

The other examples do not show Echo options for two reasons: (1) for simplicity and (2) because they lack the context for any example values to have meaning.

6. RD Lookup

To discover the resources registered with the RD, a lookup interface must be provided. This lookup interface is defined as a default, and it is assumed that RDs may also support lookups to return resource descriptions in alternative formats (e.g., JSON or CBOR link format [CORE-LINKS-JSON]) or use more advanced interfaces (e.g., supporting context- or semantic-based lookup) on different resources that are discovered independently.

RD lookup allows lookups for endpoints and resources using attributes defined in this document and for use with the CoRE Link Format. The result of a lookup request is the list of links (if any) corresponding to the type of lookup. Thus, an endpoint lookup **MUST** return a list of endpoints, and a resource lookup **MUST** return a list of links to resources.

The lookup type implemented by a lookup resource is indicated by a resource type, as per Table 1:

Lookup Type	Resource Type	Mandatory
Resource	core.rd-lookup-res	Mandatory
Endpoint	core.rd-lookup-ep	Mandatory

Table 1: Lookup Types

6.1. Resource Lookup

Resource lookup results in links that are semantically equivalent to the links submitted to the RD by the registrant. The links and link parameters returned by the lookup are equal to the originally submitted ones, except that the target reference is fully resolved and that the anchor reference is fully resolved if it is present in the lookup result at all.

Links that did not have an anchor attribute in the registration are returned without an anchor attribute. Links of which href or anchor was submitted as a (full) URI are returned with the respective attribute unmodified.

The above rules allow the client to interpret the response as links without any further knowledge of the storage conventions of the RD. The RD **MAY** replace the registration base URIs with a configured intermediate proxy, e.g., in the case of an HTTP lookup interface for CoAP endpoints.

If the base URI of a registration contains a link-local address, the RD **MUST NOT** show its links unless the lookup was made from the link on which the registered endpoint can be reached. The RD **MUST NOT** include zone identifiers in the resolved URIs.

6.2. Lookup Filtering

Using the Accept option, the requester can control whether the returned list is returned in CoRE Link Format (`application/link-format`, default) or in alternate content formats (e.g., from [CORE-LINKS-JSON]).

Multiple search criteria **MAY** be included in a lookup. All included criteria **MUST** match for a link to be returned. The RD **MUST** support matching with multiple search criteria.

A link matches a search criterion if it has an attribute of the same name and the same value, allowing for a trailing "*" wildcard operator, as in Section 4.1 of [RFC6690]. Attributes that are defined as `relation-types` (in the link-format ABNF) match if the search value matches any of their values (see Section 4.1 of [RFC6690]); for example, `?if=tag:example.net,2020:sensor` matches `;if="example.regname tag:example.net,2020:sensor";`. A resource link also matches a search criterion if its endpoint would match the criterion, and vice versa, an endpoint link matches a search criterion if any of its resource links matches it.

Note that `href` is a valid search criterion and matches target references. Like all search criteria, on a resource lookup, it can match the target reference of the resource link itself but also the registration resource of the endpoint that registered it. Queries for resource link targets **MUST** be in URI form (i.e., not relative references) and are matched against a resolved link target. Queries for endpoints **SHOULD** be expressed in path-absolute form if possible and **MUST** be expressed in URI form otherwise; the RD **SHOULD** recognize either. The `anchor` attribute is usable for resource lookups and, if queried, **MUST** be in URI form as well.

Additional query parameters "page" and "count" are used to obtain lookup results in specified increments using pagination, where count specifies how many links to return and page specifies which subset of links organized in sequential pages, each containing 'count' links, starting with

link zero and page zero. Thus, specifying a count of 10 and page of 0 will return the first 10 links in the result set (links 0-9). Specifying a count of 10 and page of 1 will return the next 'page' containing links 10-19, and so on. Unlike block-wise transfer of a complete result set, these parameters ensure that each chunk of results can be interpreted on its own. This simplifies the processing but can result in duplicate or missed items when coinciding with changes from the registration interface.

Endpoints that are interested in a lookup result repeatedly or continuously can use mechanisms such as ETag caching, resource observation [RFC7641], or any future mechanism that might allow more efficient observations of collections. These are advertised, detected, and used according to their own specifications and can be used with the lookup interface as with any other resource.

When resource observation is used, every time the set of matching links changes or the content of a matching link changes, the RD sends a notification with the matching link set. The notification contains the successful current response to the given request, especially with respect to representing zero matching links (see "Success" item below).

The lookup interface is specified as follows:

Interaction: Client -> RD

Method: GET

URI Template: {+type-lookup-location}{?page,count,search*}

URI Template Variables:

    type-lookup-location :=   RD lookup URI for a given lookup type (mandatory). The address is discovered as described in Section 4.3.

    search :=   Search criteria for limiting the number of results (optional). The search criteria are an associative array, expressed in a form-style query, as per the URI Template (see [RFC6570], Sections 2.4.2 and 3.2.8).

    page :=   Page (optional). This parameter cannot be used without the count parameter. Results are returned from result set in pages that contain 'count' links starting from index (page * count). Page numbering starts with zero.

    count :=   Count (optional). The number of results is limited to this parameter value. If the page parameter is also present, the response **MUST** only include 'count' links starting with the (page * count) link in the result set from the query. If the count parameter is not present, then the response **MUST** return all matching links in the result set. Link numbering starts with zero.

Accept: absent, `application/link-format`, or any other indicated media type representing web links

The following responses codes are defined for this interface:

Success: 2.05 (Content) or 200 (OK) with an `application/link-format` or other web link payload containing matching entries for the lookup.

The payload can contain zero links (which is an empty payload in the link format described in [RFC6690] but could also be [ ] in JSON-based formats), indicating that no entities matched the request.

6.3. Resource Lookup Examples

The examples in this section assume the existence of CoAP hosts with a default CoAP port 61616. HTTP hosts are possible and do not change the nature of the examples.

The following example shows a client performing a resource lookup with the example resource lookup locations discovered in Figure 5:

```
Req: GET /rd-lookup/res?rt=tag:example.org,2020:temperature

Res: 2.05 Content
Payload:
<coap://[2001:db8:3::123]:61616/temp>;
     rt="tag:example.org,2020:temperature"
```

Figure 19: Example of a Resource Lookup

A client that wants to be notified of new resources as they show up can use this observation:

```
Req: GET /rd-lookup/res?rt=tag:example.org,2020:light
Observe: 0

Res: 2.05 Content
Observe: 23
Payload: empty

(at a later point in time)

Res: 2.05 Content
Observe: 24
Payload:
<coap://[2001:db8:3::124]/west>;rt="tag:example.org,2020:light",
<coap://[2001:db8:3::124]/south>;rt="tag:example.org,2020:light",
<coap://[2001:db8:3::124]/east>;rt="tag:example.org,2020:light"
```

Figure 20: Example of an Observing Resource Lookup

The following example shows a client performing a paginated resource lookup:

```
Req: GET /rd-lookup/res?page=0&count=5

Res: 2.05 Content
Payload:
<coap://[2001:db8:3::123]:61616/res/0>;ct=60,
<coap://[2001:db8:3::123]:61616/res/1>;ct=60,
<coap://[2001:db8:3::123]:61616/res/2>;ct=60,
<coap://[2001:db8:3::123]:61616/res/3>;ct=60,
<coap://[2001:db8:3::123]:61616/res/4>;ct=60

Req: GET /rd-lookup/res?page=1&count=5

Res: 2.05 Content
Payload:
<coap://[2001:db8:3::123]:61616/res/5>;ct=60,
<coap://[2001:db8:3::123]:61616/res/6>;ct=60,
<coap://[2001:db8:3::123]:61616/res/7>;ct=60,
<coap://[2001:db8:3::123]:61616/res/8>;ct=60,
<coap://[2001:db8:3::123]:61616/res/9>;ct=60
```

Figure 21: Example of Paginated Resource Lookup

The following example shows a client performing a lookup of all resources of all endpoints of a given endpoint type. It assumes that two endpoints (with endpoint names `sensor1` and `sensor2`) have previously registered with their respective addresses (`coap://sensor1.example.com` and `coap://sensor2.example.com`) and posted the very payload of the 6th response of Section 5 of [RFC6690].

It demonstrates how absolute link targets stay unmodified, while relative ones are resolved:

```
Req: GET /rd-lookup/res?et=tag:example.com,2020:platform

Res: 2.05 Content
Payload:
<coap://sensor1.example.com/sensors>;ct=40;title="Sensor Index",
<coap://sensor1.example.com/sensors/temp>;rt=temperature-c;if=sensor,
<coap://sensor1.example.com/sensors/light>;rt=light-lux;if=sensor,
<http://www.example.com/sensors/t123>;rel=describedby;
    anchor="coap://sensor1.example.com/sensors/temp",
<coap://sensor1.example.com/t>;rel=alternate;
    anchor="coap://sensor1.example.com/sensors/temp",
<coap://sensor2.example.com/sensors>;ct=40;title="Sensor Index",
<coap://sensor2.example.com/sensors/temp>;rt=temperature-c;if=sensor,
<coap://sensor2.example.com/sensors/light>;rt=light-lux;if=sensor,
<http://www.example.com/sensors/t123>;rel=describedby;
    anchor="coap://sensor2.example.com/sensors/temp",
<coap://sensor2.example.com/t>;rel=alternate;
    anchor="coap://sensor2.example.com/sensors/temp"
```

Figure 22: Example of a Resource Lookup from Multiple Endpoints

6.4. Endpoint Lookup

The endpoint lookup returns links to and information about registration resources, which themselves can only be manipulated by the registering endpoint.

Endpoint registration resources are annotated with their endpoint names (ep), sectors (d, if present), and registration base URI (base; reports the registrant-EP's address if no explicit base was given), as well as a constant resource type (rt="core.rd-ep"); the lifetime (lt) is not reported. Additional endpoint attributes are added as target attributes to their endpoint link unless their specification says otherwise.

Links to endpoints **SHOULD** be presented in path-absolute form or, if required, as (full) URIs. (This ensures that the output conforms to Limited Link Format, as described in Appendix C.)

Base addresses that contain link-local addresses **MUST NOT** include zone identifiers, and such registrations **MUST NOT** be shown unless the lookup was made from the same link from which the registration was made.

While the endpoint lookup does expose the registration resources, the RD does not need to make them accessible to clients. Clients **SHOULD NOT** attempt to dereference or manipulate them.

An RD can report registrations in lookup whose URI scheme and authority differ from that of the lookup resource. Lookup clients **MUST** be prepared to see arbitrary URIs as registration resources in the results and treat them as opaque identifiers; the precise semantics of such links are left to future specifications.

The following example shows a client performing an endpoint lookup that is limited to endpoints of endpoint type `tag:example.com,2020:platform`:

```
Req: GET /rd-lookup/ep?et=tag:example.com,2020:platform

Res: 2.05 Content
Payload:
</rd/1234>;base="coap://[2001:db8:3::127]:61616";ep=node5;
    et="tag:example.com,2020:platform";ct=40;rt=core.rd-ep,
</rd/4521>;base="coap://[2001:db8:3::129]:61616";ep=node7;
    et="tag:example.com,2020:platform";ct=40;d=floor-3;
    rt=core.rd-ep
```

Figure 23: Example of Endpoint Lookup

7. Security Policies

The security policies that are applicable to an RD strongly depend on the application and are not set out normatively here.

This section provides a list of aspects that applications should consider when describing their use of the RD, without claiming to cover all cases. It uses terminology of [ACE-OAUTH-AUTHZ], in which the RD acts as the Resource Server (RS), and both registrant-EPs and lookup clients act as Clients (C) with support from an Authorization Server (AS), without the intention of ruling out other schemes (e.g., those based on certificates/Public Key Infrastructures (PKIs)).

Any, all, or none of the below can apply to an application. Which are relevant depends on its protection objectives.

Security policies are set by configuration of the RD or by choice of the implementation. Lookup clients (and, where relevant, endpoints) can only trust an RD to uphold them if it is authenticated and authorized to serve as an RD according to the application's requirements.

7.1. Endpoint Name

Whenever an RD needs to provide trustworthy results to clients doing endpoint lookup or resource lookup with filtering on the endpoint name, the RD must ensure that the registrant is authorized to use the given endpoint name. This applies both to registration and later to operations on the registration resource. It is immaterial whether the client is the registrant-EP itself or a CT is doing the registration. The RD cannot tell the difference, and CTs may use authorization credentials authorizing only operations on that particular endpoint name or a wider range of endpoint names.

It is up to the concrete security policy to describe how the endpoint name and sector are transported when certificates are used. For example, it may describe how SubjectAltName dNSName entries are mapped to endpoint and domain names.

7.1.1. Random Endpoint Names

Conversely, in applications where the RD does not check the endpoint name, the authorized registering endpoint can generate a random number (or string) that identifies the endpoint. The RD should then remember unique properties of the registrant, associate them with the registration for as long as its registration resource is active (which may be longer than the registration's lifetime), and require the same properties for operations on the registration resource.

Registrants that are prepared to pick a different identifier when their initial attempt (or attempts, in the unlikely case of two subsequent collisions) at registration is unauthorized should pick an identifier at least twice as long as would be needed to enumerate the expected number of registrants; registrants without any such recovery options should pick significantly longer endpoint names (e.g., using Universally Unique Identifier (UUID) URNs [RFC4122]).

7.2. Entered Links

When lookup clients expect that certain types of links can only originate from certain endpoints, then the RD needs to apply filtering to the links an endpoint may register.

For example, if clients use an RD to find a server that provides firmware updates, then any registrant that wants to register (or update) links to firmware sources will need to provide suitable credentials to do so, independently of its endpoint name.

Note that the impact of having undesirable links in the RD depends on the application. If the client requires the firmware server to present credentials as a firmware server, a fraudulent link's impact is limited to the client revealing its intention to obtain updates and slowing down the client until it finds a legitimate firmware server; if the client accepts any credentials from the server as long as they fit the provided URI, the impact is larger.

An RD may also require that links are only registered if the registrant is authorized to publish information about the anchor (or even target) of the link. One way to do this is to demand that the registrant present the same credentials in its role as a registering client that it would need to present in its role as a server when contacted at the resources' URI. These credentials may include using the address and port that are part of the URI. Such a restriction places severe practical limitations on the links that can be registered.

As above, the impact of undesirable links depends on the extent to which the lookup client relies on the RD. To avoid the limitations, RD applications should consider prescribing that lookup clients only use the discovered information as hints and describe which pieces of information need to be verified because they impact the application's security. A straightforward way to verify such information is to request it again from an authorized server, typically the one that hosts the target resource. That is similar to what happens in Section 4.3 when the "URI discovery" step is repeated.

7.3. Link Confidentiality

When registrants publish information in the RD that is not available to any client that would query the registrant's `/.well-known/core` interface, or when lookups to that interface are subject to stricter firewalling than lookups to the RD, the RD may need to limit which lookup clients may access the information.

In this case, the endpoint (and not the lookup clients) needs to be careful to check the RD's authorization. The RD needs to check any lookup client's authorization before revealing information directly (in resource lookup) or indirectly (when using it to satisfy a resource lookup search criterion).

7.4. Segmentation

Within a single RD, different security policies can apply.

One example of this are multi-tenant deployments separated by the sector (d) parameter. Some sectors might apply limitations on the endpoint names available, while others use a random identifier approach to endpoint names and place limits on the entered links based on their attributes instead.

Care must be taken in such setups to determine the applicable access control measures to each operation. One easy way to do that is to mandate the use of the sector parameter on all operations, as no credentials are suitable for operations across sector borders anyway.

7.5. "First Come First Remembered": A Default Policy

The "First Come First Remembered" policy is provided both as a reference example for a security policy definition and as a policy that implementations may choose to use as default policy in the absence of any other configuration. It is designed to enable efficient discovery operations even in ad hoc settings.

Under this policy, the RD accepts registrations for any endpoint name that is not assigned to an active registration resource and only accepts registration updates from the same endpoint. The policy is minimal in that it does not make any promises to lookup clients about the claims of Sections 7.2 and 7.3, and promises about the claims in Section 7.1 are limited to the lifetime of that endpoint's registration. It does however promise the endpoint that, for the duration of its registration, its links will be discoverable on the RD.

When a registration or operation is attempted, the RD **MUST** determine the client's subject name or public key:

- If the client's credentials indicate any subject name that is certified by any authority that the RD recognizes (which may be the system's trust anchor store), all such subject names are stored. With credentials based on CWT or JWT (as common with Authentication and Authorization for Constrained Environments (ACE)), the Subject (sub) claim is stored as a single name, if it exists. With X.509 certificates, the Common Name (CN) and the complete list of SubjectAltName entries are stored. In both cases, the authority that certified the claim is stored along with the subject, as the latter may only be locally unique.
- Otherwise, if the client proves possession of a private key, the matching public key is stored. This applies both to raw public keys and to the public keys indicated in certificates that failed the above authority check.
- If neither is present, a reference to the security session itself is stored. With (D)TLS, that is the connection itself or the session resumption information, if available. With OSCORE, that is the security context.

As part of the registration operation, that information is stored along with the registration resource.

The RD **MUST** accept all registrations whose registration resource is not already active, as long as they are made using a security layer supported by the RD.

Any operation on a registration resource, including registrations that lead to an existing registration resource, **MUST** be rejected by the RD unless all the stored information is found in the new request's credentials.

Note that, even though subject names are compared in this policy, they are never directly compared to endpoint names, and an endpoint cannot expect to "own" any particular endpoint name outside of an active registration -- even if a certificate says so. It is an accepted shortcoming of this approach that the endpoint has no indication of whether the RD remembers it by its subject name or public key; recognition by subject happens on a best-effort basis (given the RD may not recognize any authority). Clients **MUST** be prepared to pick a different endpoint name when rejected by the RD initially or after a change in their credentials; picking an endpoint name, as per Section 7.1.1, is an easy option for that.

For this policy to be usable without configuration, clients should not set a sector name in their registrations. An RD can set a default sector name for registrations accepted under this policy, which is especially useful in a segmented setup where different policies apply to different sectors. The configuration of such a behavior, as well as any other configuration applicable to such an RD (i.e., the set of recognized authorities), is out of scope for this document.

8.   Security Considerations

The security considerations as described in Section 5 of [RFC8288] and Section 6 of [RFC6690] apply. The `/.well-known/core` resource may be protected, e.g., using DTLS when hosted on a CoAP server, as described in [RFC7252].

Access that is limited or affects sensitive data **SHOULD** be protected, e.g., using (D)TLS or OSCORE [RFC8613]; which aspects of the RD this affects depends on the security policies of the application (see Section 7).

8.1.   Discovery

Most steps in discovery of the RD, and possibly its resources, are not covered by CoAP's security mechanisms. This will not endanger the security properties of the registrations and lookup itself (where the client requires authorization of the RD if it expects any security properties of the operation) but may leak the client's intention to third parties and allow them to slow down the process.

To mitigate that, clients can retain the RD's address, use secure discovery options (such as configured addresses), and send queries for RDs in a very general form (e.g., `?rt=core.rd*` rather than `?rt=core.rd-lookup-ep`).

8.2.   Endpoint Identification and Authentication

An endpoint (name, sector) pair is unique within the set of endpoints registered by the RD. An endpoint **MUST NOT** be identified by its protocol, port, or IP address, as these may change over the lifetime of an endpoint.

Every operation performed by an endpoint on an RD **SHOULD** be mutually authenticated using a pre-shared key, a raw public key, or certificate-based security.

Consider the following threat: two devices, A and B, are registered at a single server. Both devices have unique, per-device credentials for use with DTLS to make sure that only parties with authorization to access A or B can do so.

Now, imagine that a malicious device A wants to sabotage the device B. It uses its credentials during the DTLS exchange. Then, it specifies the endpoint name of device B as the name of its own endpoint in device A. If the server does not check whether the identifier provided in the DTLS handshake matches the identifier used at the CoAP layer, then it may be inclined to use the endpoint name for looking up what information to provision to the malicious device.

Endpoint authorization needs to be checked on registration and registration resource operations independently of whether there are configured requirements on the credentials for a given endpoint name and sector (Section 7.1) or whether arbitrary names are accepted (Section 7.1.1).

Simple registration could be used to circumvent address-based access control. An attacker would send a simple registration request with the victim's address as the source address and later look up the victim's `/.well-known/core` content in the RD. Mitigation for this is recommended in Section 5.1.

The registration resource path is visible to any client that is allowed endpoint lookup and can be extracted by resource lookup clients as well. The same goes for registration attributes that are shown as target attributes or lookup attributes. The RD needs to consider this in the choice of registration resource paths, as do administrators or endpoints in their choice of attributes.

8.3. Access Control

Access control **SHOULD** be performed separately for the RD registration and lookup API paths, as different endpoints may be authorized to register with an RD from those authorized to look up endpoints from the RD. Such access control **SHOULD** be performed in as fine-grained a level as possible. For example, access control for lookups could be performed either at the sector, endpoint, or resource level.

The precise access controls necessary (and the consequences of failure to enforce them) depend on the protection objectives of the application and the security policies (Section 7) derived from them.

8.4. Denial-of-Service Attacks

Services that run over UDP unprotected are vulnerable to unknowingly amplify and distribute a DoS attack, as UDP does not require a return routability check. Since RD lookup responses can be significantly larger than requests, RDs are prone to this.

[RFC7252] describes this at length in its Section 11.3, including some mitigation by using small block sizes in responses. [RFC9175] updates that by describing a source address verification mechanism using the Echo option.

8.5. Skipping Freshness Checks

When RD-based applications are built in which request freshness checks are not performed, these concerns need to be balanced:

- When alterations to registration attributes are reordered, an attacker may create any combination of attributes ever set, with the attack difficulty determined by the security layer's replay properties.

  For example, if Figure 18 were conducted without freshness assurances, an attacker could later reset the lifetime back to 7200. Thus, the device is made unreachable to lookup clients.

- When registration updates without query parameters (which just serve to restart the lifetime) can be reordered, an attacker can use intercepted messages to give the appearance of the device being alive to the RD.

  This is unacceptable when the RD's security policy promises reachability of endpoints (e.g., when disappearing devices would trigger further investigation) but may be acceptable with other policies.

9. IANA Considerations

9.1. Resource Types

IANA has added the following values to the "Resource Type (rt=) Link Target Attribute Values" subregistry of the "Constrained RESTful Environments (CoRE) Parameters" registry defined in [RFC6690]:

Value	Description	Reference
core.rd	Directory resource of an RD	RFC 9176, Section 4.3
core.rd-lookup-res	Resource lookup of an RD	RFC 9176, Section 4.3
core.rd-lookup-ep	Endpoint lookup of an RD	RFC 9176, Section 4.3
core.rd-ep	Endpoint resource of an RD	RFC 9176, Section 6

Table 2: Additions to Resource Type (rt=) Link Target Attribute Values Subregistry

9.2. IPv6 ND Resource Directory Address Option

IANA has registered one new ND option type in the "IPv6 Neighbor Discovery Option Formats" subregistry of the "Internet Control Message Protocol version 6 (ICMPv6) Parameters" registry:

Type	Description	Reference
41	Resource Directory Address Option	RFC 9176

Table 3: Addition to IPv6 Neighbor Discovery Option Formats Subregistry

9.3. RD Parameters Registry

This specification defines a new subregistry for registration and lookup parameters called "RD Parameters" within the "Constrained RESTful Environments (CoRE) Parameters" registry. Although this specification defines a basic set of parameters, it is expected that other standards that make use of this interface will define new ones.

Each entry in the registry must include:

- the human-readable name of the parameter,
- the short name, as used in query parameters or target attributes,
- syntax and validity requirements (if any),
- indication of whether it can be passed as a query parameter at registration of endpoints, passed as a query parameter in lookups, or expressed as a target attribute,
- a description, and
- a link to reference documentation.

The query parameter **MUST** be both a valid URI query key [RFC3986] and a token as used in [RFC8288].

The reference documentation must give details on whether the parameter can be updated and how it is to be processed in lookups.

The mechanisms around new RD parameters should be designed in such a way that they tolerate RD implementations that are unaware of the parameter and expose any parameter passed at registration or updates in endpoint lookups. (For example, if a parameter used at registration were to be confidential, the registering endpoint should be instructed to only set that parameter if the RD advertises support for keeping it confidential at the discovery step.)

Initial entries in this subregistry are as follows:

Name	Short	Validity	Use	Description
Endpoint Name	ep	Unicode*	RLA	Name of the endpoint
Lifetime	lt	1-4294967295	R	Lifetime of the registration in seconds
Sector	d	Unicode*	RLA	Sector to which this endpoint belongs

Name	Short	Validity	Use	Description
Registration Base URI	base	URI	RLA	The scheme, address, port, and path at which this server is available
Page	page	Integer	L	Used for pagination
Count	count	Integer	L	Used for pagination
Endpoint Type	et	RFC 9176, Section 9.3.1	RLA	Semantic type of the endpoint (see RFC 9176, Section 9.4)

Table 4: New RD Parameters Registry

Where:

Short: Short name used in query parameters or target attributes

Validity:

Unicode* =   up to 63 bytes of UTF-8-encoded Unicode, with no control characters as per Section 5

Use:

R =   used at registration
L =   used at lookup
A =   expressed in the target attribute

The descriptions for the options defined in this document are only summarized here. To which registrations they apply and when they are to be shown are described in the respective sections of this document. All their reference documentation entries point to this document.

The IANA policy for future additions to the subregistry is Expert Review, as described in [RFC8126]. The evaluation should consider formal criteria, duplication of functionality (i.e., is the new entry redundant with an existing one?), topical suitability (e.g., is the described property actually a property of the endpoint and not a property of a particular resource, in which case it should go into the payload of the registration and need not be registered?), and the potential for conflict with commonly used target attributes (e.g., `if` could be used as a parameter for conditional registration if it were not to be used in lookup or attributes but would make a bad parameter for lookup because a resource lookup with an `if` query parameter could ambiguously filter by the registered endpoint property or the target attribute [RFC6690]).

9.3.1. Full Description of the "Endpoint Type" RD Parameter

An endpoint registering at an RD can describe itself with endpoint types, similar to how resources are described with resource types in [RFC6690]. An endpoint type is expressed as a string, which can be either a URI or one of the values defined in the "Endpoint Type (et=) RD Parameter Values" subregistry. Endpoint types can be passed in the `et` query parameter as part

of extra-attrs at the "registration" step of Section 5, are shown on endpoint lookups using the `et` target attribute, and can be filtered for using `et` as a search criterion in resource and endpoint lookup. Multiple endpoint types are given as separate query parameters or link attributes.

Note that the endpoint type differs from the resource type in that it uses multiple attributes rather than space-separated values. As a result, RDs implementing this specification automatically support correct filtering in the lookup interfaces from the rules for unknown endpoint attributes.

9.4. Endpoint Type (et=) RD Parameter Values

This specification establishes a new subregistry called "Endpoint Type (et=) RD Parameter Values" within the "Constrained RESTful Environments (CoRE) Parameters" registry. The registry properties (required policy, requirements, and template) are identical to those of the "Resource Type (rt=) Link Target Attribute Values" subregistry defined in [RFC6690]; in short, the review policy is IETF Review for values starting with "core" and Specification Required for others.

The requirements to be enforced are:

- The values **MUST** be related to the purpose described in Section 9.3.1.
- The registered values **MUST** conform to the ABNF reg-rel-type definition of [RFC6690] and **MUST NOT** be a URI.
- It is recommended to use the period "." character for segmentation.

The initial contents of the registry are as follows:

Value	Description	Reference
core.rd-group	An application group, as described in RFC 9176, Appendix A.	RFC 9176

Table 5: New Endpoint Type (et=) RD Parameter Values Registry

9.5. Multicast Address Registration

IANA has assigned the following multicast addresses for use by CoAP nodes:

IPv4 -- "All CoRE Resource Directories" address 224.0.1.190, in the "Internetwork Control Block (224.0.1.0 - 224.0.1.255 (224.0.1/24))" subregistry within the "IPv4 Multicast Address Space Registry". As the address is used for discovery that may span beyond a single network, it has come from the Internetwork Control Block (224.0.1.x) [RFC5771].

IPv6 -- "All CoRE Resource Directories" address ff0x::fe, in the "Variable Scope Multicast Addresses" subregistry within the "IPv6 Multicast Address Space Registry" [RFC3307]. Note that there is a distinct multicast address for each scope that interested CoAP nodes should listen to; CoAP needs the link-local and site-local scopes only.

9.6. Well-Known URIs

IANA has registered the URI suffix "rd" in the "Well-Known URIs" registry as follows:

URI Suffix	Change Controller	Reference	Status
rd	IETF	RFC 9176	permanent

Table 6: Addition to Well-Known URIs Registry

9.7. Service Name and Transport Protocol Port Number Registry

IANA has added four new items to the "Service Name and Transport Protocol Port Number Registry" as follows:

Service Name	Transport Protocol	Description	Reference
core-rd	udp	Resource Directory accessed using CoAP	RFC 9176
core-rd-dtls	udp	Resource Directory accessed using CoAP over DTLS	RFC 9176
core-rd	tcp	Resource Directory accessed using CoAP over TCP	RFC 9176
core-rd-tls	tcp	Resource Directory accessed using CoAP over TLS	RFC 9176

Table 7: Additions to Service Name and Transport Protocol Port Number Registry

10. Examples

Two examples are presented: a lighting installation example in Section 10.1 and a Lightweight M2M (LwM2M) example in Section 10.2.

10.1. Lighting Installation

This example shows a simplified lighting installation that makes use of the RD with a CoAP interface to facilitate the installation and startup of the application code in the lights and sensors. In particular, the example leads to the definition of a group and the enabling of the corresponding multicast address, as described in Appendix A. No conclusions must be drawn on the realization of actual installation or naming procedures, because the example only emphasizes some of the issues that may influence the use of the RD and does not pretend to be normative.

10.1.1. Installation Characteristics

The example assumes that the installation is managed. That means that a Commissioning Tool (CT) is used to authorize the addition of nodes, name them, and name their services. The CT can be connected to the installation in many ways: the CT can be part of the installation network, connected by Wi-Fi to the installation network, connected via GPRS link, or connected by another method.

It is assumed that there are two naming authorities for the installation: (1) the network manager that is responsible for the correct operation of the network and the connected interfaces and (2) the lighting manager that is responsible for the correct functioning of networked lights and sensors. The result is the existence of two naming schemes coming from the two managing entities.

The example installation consists of one presence sensor and two luminaries, luminary1 and luminary2, each with their own wireless interface. Each luminary contains three lamps: left, right, and middle. Each luminary is accessible through one endpoint. For each lamp, a resource exists to modify the settings of a lamp in a luminary. The purpose of the installation is that the presence sensor notifies the presence of persons to a group of lamps. The group of lamps consists of the middle and left lamps of luminary1 and the right lamp of luminary2.

Before commissioning by the lighting manager, the network is installed, and access to the interfaces is proven to work by the network manager.

At the moment of installation, the network under installation is not necessarily connected to the DNS infrastructure. Therefore, Stateless Address Autoconfiguration (SLAAC) IPv6 addresses are assigned to CT, RD, luminaries, and the sensor. The addresses shown in Table 8 below stand in for these in the following examples.

Name	IPv6 address
luminary1	2001:db8:4::1
luminary2	2001:db8:4::2
Presence sensor	2001:db8:4::3
RD	2001:db8:4::ff

Table 8: Addresses Used in the Examples

In Section 10.1.2, the use of RD during installation is presented.

10.1.2. RD Entries

It is assumed that access to the DNS infrastructure is not always possible during installation. Therefore, the SLAAC addresses are used in this section.

For discovery, the resource types (rt) of the devices are important. The lamps in the luminaries have rt=tag:example.com,2020:light, and the presence sensor has rt=tag:example.com,2020:p-sensor. The endpoints have names that are relevant to the light installation manager. In this case, luminary1, luminary2, and the presence sensor are located in room 2-4-015, where luminary1 is located at the window and luminary2 and the presence sensor are located at the door. The endpoint names reflect this physical location. The middle, left, and right lamps are accessed via path /light/middle, /light/left, and /light/right, respectively. The identifiers relevant to the RD are shown in Table 9.

Name	Endpoint	Resource Path	Resource Type
luminary1	lm_R2-4-015_wndw	/light/left	tag:example.com,2020:light
luminary1	lm_R2-4-015_wndw	/light/middle	tag:example.com,2020:light
luminary1	lm_R2-4-015_wndw	/light/right	tag:example.com,2020:light
luminary2	lm_R2-4-015_door	/light/left	tag:example.com,2020:light
luminary2	lm_R2-4-015_door	/light/middle	tag:example.com,2020:light
luminary2	lm_R2-4-015_door	/light/right	tag:example.com,2020:light
Presence sensor	ps_R2-4-015_door	/ps	tag:example.com,2020:p-sensor

Table 9: RD Identifiers

It is assumed that the CT has performed RD discovery and has received a response like the one in the example in Section 4.3.

The CT inserts the endpoints of the luminaries and the sensor in the RD using the registration base URI parameter (base) to specify the interface address:

```
Req: POST coap://[2001:db8:4::ff]/rd
  ?ep=lm_R2-4-015_wndw&base=coap://[2001:db8:4::1]&d=R2-4-015
Payload:
</light/left>;rt="tag:example.com,2020:light",
</light/middle>;rt="tag:example.com,2020:light",
</light/right>;rt="tag:example.com,2020:light"

Res: 2.01 Created
Location-Path: /rd/4521

Req: POST coap://[2001:db8:4::ff]/rd
  ?ep=lm_R2-4-015_door&base=coap://[2001:db8:4::2]&d=R2-4-015
Payload:
</light/left>;rt="tag:example.com,2020:light",
</light/middle>;rt="tag:example.com,2020:light",
</light/right>;rt="tag:example.com,2020:light"

Res: 2.01 Created
Location-Path: /rd/4522

Req: POST coap://[2001:db8:4::ff]/rd
  ?ep=ps_R2-4-015_door&base=coap://[2001:db8:4::3]&d=R2-4-015
Payload:
</ps>;rt="tag:example.com,2020:p-sensor"

Res: 2.01 Created
Location-Path: /rd/4523
```

Figure 24: Example of Registrations a CT Enters into an RD

The sector name d=R2-4-015 has been added for an efficient lookup because filtering on the "ep" name is more awkward. The same sector name is communicated to the two luminaries and the presence sensor by the CT.

The group is specified in the RD. The base parameter is set to the site-local multicast address allocated to the group. In the POST in the example below, the resources supported by all group members are published.

```
Req: POST coap://[2001:db8:4::ff]/rd
  ?ep=grp_R2-4-015&et=core.rd-group&base=coap://[ff05::1]
Payload:
</light/left>;rt="tag:example.com,2020:light",
</light/middle>;rt="tag:example.com,2020:light",
</light/right>;rt="tag:example.com,2020:light"

Res: 2.01 Created
Location-Path: /rd/501
```

Figure 25: Example of a Multicast Group a CT Enters into an RD

After the filling of the RD by the CT, the application in the luminaries can learn to which groups they belong and enable their interface for the multicast address.

The luminary, knowing its sector and being configured to join any group containing lights, searches for candidate groups and joins them:

```
Req: GET coap://[2001:db8:4::ff]/rd-lookup/ep
  ?d=R2-4-015&et=core.rd-group&rt=light

Res: 2.05 Content
Payload:
</rd/501>;ep=grp_R2-4-015;et=core.rd-group;
         base="coap://[ff05::1]";rt=core.rd-ep
```

Figure 26: Example of a Lookup Exchange to Find Suitable Multicast Addresses

From the returned base parameter value, the luminary learns the multicast address of the multicast group.

The presence sensor can learn the presence of groups that support resources with rt=tag:example.com,2020:light in its own sector by sending the same request, as used by the luminary. The presence sensor learns the multicast address to use for sending messages to the luminaries.

10.2. OMA Lightweight M2M (LwM2M)

OMA LwM2M is a profile for device services based on CoAP, providing interfaces and operations for device management and device service enablement.

An LwM2M server is an instance of an LwM2M middleware service layer, containing an RD ([LwM2M], starting at page 36).

That RD only implements the registration interface, and no lookup is implemented. Instead, the LwM2M server provides access to the registered resources in a similar way to a reverse proxy.

The location of the LwM2M server and RD URI path is provided by the LwM2M bootstrap process, so no dynamic discovery of the RD is used. LwM2M servers and endpoints are not required to implement the `/.well-known/core` resource.

11. References

11.1. Normative References

[RFC2119]   Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <https://www.rfc-editor.org/info/rfc2119>.

[RFC3986]   Berners-Lee, T., Fielding, R., and L. Masinter, "Uniform Resource Identifier (URI): Generic Syntax", STD 66, RFC 3986, DOI 10.17487/RFC3986, January 2005, <https://www.rfc-editor.org/info/rfc3986>.

[RFC6570]   Gregorio, J., Fielding, R., Hadley, M., Nottingham, M., and D. Orchard, "URI Template", RFC 6570, DOI 10.17487/RFC6570, March 2012, <https://www.rfc-editor.org/info/rfc6570>.

[RFC6690]   Shelby, Z., "Constrained RESTful Environments (CoRE) Link Format", RFC 6690, DOI 10.17487/RFC6690, August 2012, <https://www.rfc-editor.org/info/rfc6690>.

[RFC6763]   Cheshire, S. and M. Krochmal, "DNS-Based Service Discovery", RFC 6763, DOI 10.17487/RFC6763, February 2013, <https://www.rfc-editor.org/info/rfc6763>.

[RFC7230]   Fielding, R., Ed. and J. Reschke, Ed., "Hypertext Transfer Protocol (HTTP/1.1): Message Syntax and Routing", RFC 7230, DOI 10.17487/RFC7230, June 2014, <https://www.rfc-editor.org/info/rfc7230>.

[RFC7252]   Shelby, Z., Hartke, K., and C. Bormann, "The Constrained Application Protocol (CoAP)", RFC 7252, DOI 10.17487/RFC7252, June 2014, <https://www.rfc-editor.org/info/rfc7252>.

[RFC8126]   Cotton, M., Leiba, B., and T. Narten, "Guidelines for Writing an IANA Considerations Section in RFCs", BCP 26, RFC 8126, DOI 10.17487/RFC8126, June 2017, <https://www.rfc-editor.org/info/rfc8126>.

[RFC8174]   Leiba, B., "Ambiguity of Uppercase vs Lowercase in RFC 2119 Key Words", BCP 14, RFC 8174, DOI 10.17487/RFC8174, May 2017, <https://www.rfc-editor.org/info/rfc8174>.

[RFC8288]   Nottingham, M., "Web Linking", RFC 8288, DOI 10.17487/RFC8288, October 2017, <https://www.rfc-editor.org/info/rfc8288>.

[RFC9175]   Amsüss, C., Preuß Mattsson, J., and G. Selander, "Constrained Application Protocol (CoAP): Echo, Request-Tag, and Token Processing", RFC 9175, DOI 10.17487/RFC9175, February 2022, <https://www.rfc-editor.org/info/rfc9175>.

11.2. Informative References

[ACE-OAUTH-AUTHZ]   Seitz, L., Selander, G., Wahlstroem, E., Erdtman, S., and H. Tschofenig, "Authentication and Authorization for Constrained Environments (ACE) using the OAuth 2.0 Framework (ACE-OAuth)", Work in Progress, Internet-Draft, draft-ietf-ace-oauth-authz-46, 8 November 2021, <https://datatracker.ietf.org/doc/html/draft-ietf-ace-oauth-authz-46>.

[COAP-PROT-NEG]   Silverajan, B. and M. Ocak, "CoAP Protocol Negotiation", Work in Progress, Internet-Draft, draft-silverajan-core-coap-protocol-negotiation-09, 2 July 2018, <https://datatracker.ietf.org/doc/html/draft-silverajan-core-coap-protocol-negotiation-09>.

[CORE-CORAL]   Amsüss, C. and T. Fossati, "The Constrained RESTful Application Language (CoRAL)", Work in Progress, Internet-Draft, draft-ietf-core-coral-05, 7 March 2022, <https://datatracker.ietf.org/doc/html/draft-ietf-core-coral-05>.

[CORE-LINKS-JSON]   Li, K., Rahman, A., and C. Bormann, Ed., "Representing Constrained
RESTful Environments (CoRE) Link Format in JSON and CBOR", Work in
Progress, Internet-Draft, draft-ietf-core-links-json-10, 26 February 2018, <https://
datatracker.ietf.org/doc/html/draft-ietf-core-links-json-10>.

[CORE-RD-DNS-SD]   van der Stok, P., Koster, M., and C. Amsuess, "CoRE Resource Directory:
DNS-SD mapping", Work in Progress, Internet-Draft, draft-ietf-core-rd-dns-sd-05,
7 July 2019, <https://datatracker.ietf.org/doc/html/draft-ietf-core-rd-dns-sd-05>.

[ER]   Chen, P., "The entity-relationship model--toward a unified view of data", ACM
Transactions on Database Systems, Vol. 1, pp. 9-36, DOI 10.1145/320434.320440,
March 1976, <https://doi.org/10.1145/320434.320440>.

[LwM2M]   Open Mobile Alliance, "Lightweight Machine to Machine Technical Specification:
Transport Bindings (Candidate Version 1.1)", June 2018, <https://
openmobilealliance.org/RELEASE/LightweightM2M/V1_1-20180612-C/OMA-TS-
LightweightM2M_Transport-V1_1-20180612-C.pdf>.

[RFC3306]   Haberman, B. and D. Thaler, "Unicast-Prefix-based IPv6 Multicast Addresses",
RFC 3306, DOI 10.17487/RFC3306, August 2002, <https://www.rfc-editor.org/info/
rfc3306>.

[RFC3307]   Haberman, B., "Allocation Guidelines for IPv6 Multicast Addresses", RFC 3307,
DOI 10.17487/RFC3307, August 2002, <https://www.rfc-editor.org/info/rfc3307>.

[RFC3849]   Huston, G., Lord, A., and P. Smith, "IPv6 Address Prefix Reserved for
Documentation", RFC 3849, DOI 10.17487/RFC3849, July 2004, <https://www.rfc-
editor.org/info/rfc3849>.

[RFC4122]   Leach, P., Mealling, M., and R. Salz, "A Universally Unique IDentifier (UUID) URN
Namespace", RFC 4122, DOI 10.17487/RFC4122, July 2005, <https://www.rfc-
editor.org/info/rfc4122>.

[RFC4944]   Montenegro, G., Kushalnagar, N., Hui, J., and D. Culler, "Transmission of IPv6
Packets over IEEE 802.15.4 Networks", RFC 4944, DOI 10.17487/RFC4944,
September 2007, <https://www.rfc-editor.org/info/rfc4944>.

[RFC5771]   Cotton, M., Vegoda, L., and D. Meyer, "IANA Guidelines for IPv4 Multicast
Address Assignments", BCP 51, RFC 5771, DOI 10.17487/RFC5771, March 2010,
<https://www.rfc-editor.org/info/rfc5771>.

[RFC6724]   Thaler, D., Ed., Draves, R., Matsumoto, A., and T. Chown, "Default Address
Selection for Internet Protocol Version 6 (IPv6)", RFC 6724, DOI 10.17487/
RFC6724, September 2012, <https://www.rfc-editor.org/info/rfc6724>.

[RFC6775]   Shelby, Z., Ed., Chakrabarti, S., Nordmark, E., and C. Bormann, "Neighbor
Discovery Optimization for IPv6 over Low-Power Wireless Personal Area
Networks (6LoWPANs)", RFC 6775, DOI 10.17487/RFC6775, November 2012,
<https://www.rfc-editor.org/info/rfc6775>.

[RFC6874]   Carpenter, B., Cheshire, S., and R. Hinden, "Representing IPv6 Zone Identifiers in
              Address Literals and Uniform Resource Identifiers", RFC 6874, DOI 10.17487/
              RFC6874, February 2013, <https://www.rfc-editor.org/info/rfc6874>.

[RFC7228]   Bormann, C., Ersue, M., and A. Keranen, "Terminology for Constrained-Node
              Networks", RFC 7228, DOI 10.17487/RFC7228, May 2014, <https://www.rfc-
              editor.org/info/rfc7228>.

[RFC7641]   Hartke, K., "Observing Resources in the Constrained Application Protocol
              (CoAP)", RFC 7641, DOI 10.17487/RFC7641, September 2015, <https://www.rfc-
              editor.org/info/rfc7641>.

[RFC8106]   Jeong, J., Park, S., Beloeil, L., and S. Madanapalli, "IPv6 Router Advertisement
              Options for DNS Configuration", RFC 8106, DOI 10.17487/RFC8106, March 2017,
              <https://www.rfc-editor.org/info/rfc8106>.

[RFC8132]   van der Stok, P., Bormann, C., and A. Sehgal, "PATCH and FETCH Methods for the
              Constrained Application Protocol (CoAP)", RFC 8132, DOI 10.17487/RFC8132,
              April 2017, <https://www.rfc-editor.org/info/rfc8132>.

[RFC8141]   Saint-Andre, P. and J. Klensin, "Uniform Resource Names (URNs)", RFC 8141, DOI
              10.17487/RFC8141, April 2017, <https://www.rfc-editor.org/info/rfc8141>.

[RFC8613]   Selander, G., Mattsson, J., Palombini, F., and L. Seitz, "Object Security for
              Constrained RESTful Environments (OSCORE)", RFC 8613, DOI 10.17487/
              RFC8613, July 2019, <https://www.rfc-editor.org/info/rfc8613>.

[T2TRG-REL-IMPL]   Bormann, C., "impl-info: A link relation type for disclosing implementation
              information", Work in Progress, Internet-Draft, draft-bormann-t2trg-rel-impl-02,
              27 September 2020, <https://datatracker.ietf.org/doc/html/draft-bormann-t2trg-
              rel-impl-02>.

Appendix A.   Groups Registration and Lookup

The RD-Group's usage pattern allows announcing application groups inside an RD.

Groups are represented by endpoint registrations. Their base address is a multicast address, and
they **SHOULD** be entered with the endpoint type `core.rd-group`. The endpoint name can also be
referred to as a group name in this context.

The registration is inserted into the RD by a Commissioning Tool, which might also be known as a
group manager here. It performs third-party registration and registration updates.

The links it registers **SHOULD** be available on all members that join the group. Depending on the
application, members that lack some resources **MAY** be permissible if requests to them fail
gracefully.

The following example shows a CT registering a group with the name "lights", which provides two resources. The directory resource path /rd is an example RD location discovered in a request similar to Figure 5. The group address in the example is constructed from the reserved 2001:db8:: prefix in [RFC3849] as a unicast-prefix-based site-local address (see [RFC3306]).

```
Req: POST coap://rd.example.com/rd?ep=lights&et=core.rd-group
                          &base=coap://[ff35:30:2001:db8:f1::8000:1]
Content-Format: 40
Payload:
</light>;rt="tag:example.com,2020:light";
     if="tag:example.net,2020:actuator",
</color-temperature>;if="tag:example.net,2020:parameter";u=K

Res: 2.01 Created
Location-Path: /rd/12
```

Figure 27: Example Registration of a Group

In this example, the group manager can easily permit devices that have no writable color-temperature to join, as they would still respond to brightness-changing commands. Had the group instead contained a single resource that sets brightness and color-temperature atomically, endpoints would need to support both properties.

The resources of a group can be looked up like any other resource, and the group registrations (along with any additional registration parameters) can be looked up using the endpoint lookup interface.

The following example shows a client performing an endpoint lookup for all groups:

```
Req: GET /rd-lookup/ep?et=core.rd-group

Res: 2.05 Content
Payload:
</rd/12>;ep=lights&et=core.rd-group;
        base="coap://[ff35:30:2001:f1:db8::8000:1]";rt=core.rd-ep
```

Figure 28: Example Lookup of Groups

The following example shows a client performing a lookup of all resources of all endpoints (groups) with et=core.rd-group:

```
Req: GET /rd-lookup/res?et=core.rd-group

Res: 2.05 Content
Payload:
<coap://[ff35:30:2001:db8:f1::8000:1]/light>;
     rt="tag:example.com,2020:light";
     if="tag:example.net,2020:actuator",
<coap://[ff35:30:2001:db8:f1::8000:1]/color-temperature>;
     if="tag:example.net,2020:parameter";u=K,
```

Figure 29: Example Lookup of Resources Inside Groups

Appendix B.   Web Links and the Resource Directory

Understanding the semantics of a link-format document and its URI references is a journey through different documents ([RFC3986] defining URIs, [RFC6690] defining link-format documents based on [RFC8288], which defines Link header fields, and [RFC7252] providing the transport). This appendix summarizes the mechanisms and semantics at play from an entry in `/.well-known/core` to a resource lookup.

This text is primarily aimed at people entering the field of Constrained Restful Environments from applications that previously did not use web mechanisms.

B.1.   A Simple Example

Let's start this example with a very simple host, `2001:db8:f0::1`. A client that follows classical CoAP discovery ([RFC7252], Section 7) sends the following multicast request to learn about neighbors supporting resources with resource-type "temperature".

The client sends a link-local multicast:

```
Req: GET coap://[ff02::fd]:5683/.well-known/core?rt=temperature

Res: 2.05 Content
Payload:
</sensors/temp>;rt=temperature;ct=0
```

Figure 30: Example of Direct Resource Discovery

where the response is sent by the server, `[2001:db8:f0::1]:5683`.

While a practical client side implementation might just go ahead and create a new request to `[2001:db8:f0::1]:5683` with Uri-Path `sensors` and `temp`, the full resolution steps for insertion into and retrieval from the RD without any shortcuts are as follows.

B.1.1. Resolving the URIs

The client parses the single returned link. Its target (sometimes called "href") is `/sensors/temp`, which is a relative URI that needs resolving. The base URI `coap://[ff02::fd]:5683/.well-known/core` is used to resolve the reference against `/sensors/temp`.

The base URI of the requested resource can be composed from the options of the CoAP GET request by following the steps of [RFC7252], Section 6.5 (with an addition at the end of Section 8.2) into `coap://[2001:db8:f0::1]/.well-known/core`.

Because `/sensors/temp` starts with a single slash, the link's target is resolved by replacing the path `/.well-known/core` from the base URI ([RFC3986], Section 5.2) with the relative target URI `/sensors/temp` into `coap://[2001:db8:f0::1]/sensors/temp`.

B.1.2. Interpreting Attributes and Relations

Some more information about the link's target can be obtained from the payload: the resource type of the target is "temperature", and its content format is text/plain (ct=0).

A relation in a web link is a three-part statement that specifies a named relation between the so-called "context resource" and the target resource, like "*This page* has *its table of contents* at */toc.html*". In link-format documents, there is an implicit "host relation" specified with default parameter rel="hosts".

In our example, the context resource of the link is implied to be `coap:://[2001:db8:f0::1]` by the default value of the anchor (see Appendix B.4). A full English expression of the "host relation" is:

> `coap://[2001:db8:f0::1]` is hosting the resource `coap://[2001:db8:f0::1]/sensors/temp`, which is of the resource type "temperature" and can be read in the text/plain content format.

B.2. A Slightly More Complex Example

Omitting the `rt=temperature` filter, the discovery query would have given some more links in the payload:

```
Req: GET coap://[ff02::fd]:5683/.well-known/core

Res: 2.05 Content
Payload:
</sensors/temp>;rt=temperature;ct=0,
</sensors/light>;rt=light-lux;ct=0,
</t>;anchor="/sensors/temp";rel=alternate,
<http://www.example.com/sensors/t123>;anchor="/sensors/temp";
    rel=describedby
```

Figure 31: Extended Example of Direct Resource Discovery

Parsing the third link, the client encounters the "anchor" parameter. It is a URI relative to the base URI of the request and is thus resolved to `coap://[2001:db8:f0::1]/sensors/temp`. That is the context resource of the link, so the "rel" statement is not about the target and the base URI any more but about the target and the resolved URI. Thus, the third link could be read as:

> `coap://[2001:db8:f0::1]/sensors/temp` has an alternate representation at `coap://[2001:db8:f0::1]/t`.

Following the same resolution steps, the fourth link can be read as `coap://[2001:db8:f0::1]/sensors/temp` is described by `http://www.example.com/sensors/t123`.

B.3. Enter the Resource Directory

The RD tries to carry the semantics obtainable by classical CoAP discovery over to the resource lookup interface as faithfully as possible.

For the following queries, we will assume that the simple host has used simple registration to register at the RD that was announced to it, sending this request from its UDP port `[2001:db8:f0::1]:6553`:

```
Req: POST coap://[2001:db8:f0::ff]/.well-known/rd?ep=simple-host1

Res: 2.04 Changed
```

Figure 32: Example of a Simple Registration

The RD would have accepted the registration and queried the simple host's `/.well-known/core` by itself. As a result, the host is registered as an endpoint in the RD with the name "simple-host1". The registration is active for 90000 seconds, and the endpoint registration base URI is `coap://[2001:db8:f0::1]`, following the resolution steps described in Appendix B.1.1. It should be remarked that the base URI constructed that way always yields a URI of the form scheme:// authority without path suffix.

If the client now queries the RD as it would previously have issued a multicast request, it would go through the RD discovery steps by fetching `coap://[2001:db8:f0::ff]/.well-known/core?rt=core.rd-lookup-res`, obtain `coap://[2001:db8:f0::ff]/rd-lookup/res` as the resource lookup endpoint, and ask it for all temperature resources:

```
Req: GET coap://[2001:db8:f0::ff]/rd-lookup/res?rt=temperature

Res: 2.05 Content
Payload:
<coap://[2001:db8:f0::1]/sensors/temp>;rt=temperature;ct=0
```

Figure 33: Example Exchange Performing Resource Lookup

This is not *literally* the same response that it would have received from a multicast request, but it contains the equivalent statement:

> `coap://[2001:db8:f0::1]` is hosting the resource `coap://[2001:db8:f0::1]/sensors/temp`, which is of the resource type "temperature" and can be accessed using the text/plain content format.

To complete the examples, the client could also query all resources hosted at the endpoint with the known endpoint name "simple-host1":

```
Req: GET coap://[2001:db8:f0::ff]/rd-lookup/res?ep=simple-host1

Res: 2.05 Content
Payload:
<coap://[2001:db8:f0::1]/sensors/temp>;rt=temperature;ct=0,
<coap://[2001:db8:f0::1]/sensors/light>;rt=light-lux;ct=0,
<coap://[2001:db8:f0::1]/t>;
    anchor="coap://[2001:db8:f0::1]/sensors/temp";rel=alternate,
<http://www.example.com/sensors/t123>;
    anchor="coap://[2001:db8:f0::1]/sensors/temp";rel=describedby
```

Figure 34: Extended Example Exchange Performing Resource Lookup

All the target and anchor references are already in absolute form there, which don't need to be resolved any further.

Had the simple host done an equivalent full registration with a base= parameter (e.g., ?ep=simple-host1&base=coap+tcp://sh1.example.com), that context would have been used to resolve the relative anchor values instead, giving the following and analogous links:

```
<coap+tcp://sh1.example.com/sensors/temp>;rt=temperature;ct=0
```

Figure 35: Example Payload of a Response to a Resource Lookup with a Dedicated Base URI

B.4. A Note on Differences between Link-Format and Link Header Fields

While link-format and Link header fields look very similar and are based on the same model of typed links, there are some differences between [RFC6690] and [RFC8288]. When implementing an RD or interacting with an RD, care must be taken to follow the behavior described in [RFC6690] whenever `application/link-format` representations are used.

- "Default value of anchor": Under both [RFC6690] and [RFC8288], relative references in the term inside the angle brackets (the target) and the anchor attribute are resolved against the relevant base URI (which usually is the URI used to retrieve the entity) and independent of each other.

  When, in a Link header [RFC8288], the anchor attribute is absent, the link's context is the URI of the selected representation (and usually equal to the base URI).

In links per [RFC6690], if the anchor attribute is absent, the default value is the Origin of (for all relevant cases, the URI reference `/` resolved against) the link's target.

- There is no percent encoding in link-format documents.

A link-format document is a UTF-8-encoded string of Unicode characters and does not have percent encoding, while Link header fields are practically ASCII strings that use percent encoding for non-ASCII characters, stating the encoding explicitly when required.

For example, while a Link header field in a page about a Swedish city might read:

```
Link: </temperature/Malm%C3%B6>;rel=live-environment-data
```

a link-format document from the same source might describe the link as:

```
</temperature/Malmö>;rel=live-environment-data
```

Appendix C.   Limited Link Format

The CoRE Link Format, as described in [RFC6690], has been interpreted differently by implementers, and a strict implementation rules out some use cases of an RD (e.g., base values with path components in combination with absent anchors).

This appendix describes a subset of link format documents called the Limited Link Format. The one rule herein is not very limiting in practice -- all examples in [RFC6690] and all deployments the authors are aware of already stick to them -- but eases the implementation of RD servers.

It is applicable to representations in the `application/link-format` media type and any other media types that inherit [RFC6690], Section 2.1.

A link format representation is in the Limited Link Format if, for each link in it, the following applies:

All URI references either follow the URI or the path-absolute ABNF rule of [RFC3986] (i.e., the target and anchor each either start with a scheme or with a single slash).

Acknowledgments

Authors' Addresses

Christian Amsüss (EDITOR)
Email: christian@amsuess.com

Zach Shelby
Edge Impulse
3031 Tisch Way
San Jose, 95128
United States of America
Email: zach@edgeimpulse.com

Michael Koster
PassiveLogic
524 H Street
Antioch, CA 94509
United States of America
Phone: +1-707-502-5136
Email: michaeljohnkoster@gmail.com

Carsten Bormann
Universität Bremen TZI
Postfach 330440
D-28359 Bremen
Germany
Phone: +49-421-218-63921
Email: cabo@tzi.org

Peter van der Stok
vanderstok consultancy
Email: stokcons@bbhmail.nl